



AI For Ardupilot Developers

Neural Networks

OlivierJB

2020 Ardupilot Developer Conference

AI Caveat emptor

- Definition of intelligence? Many! None perfect
- One simple definition of intelligence, but has issues :
The ability to acquire and apply knowledge and skills.” (Dictionary)
- Ability to acquire ... : “Machine learning”
- Lack of precise definition unfortunately contributes to misunderstandings and hype.
 - What's AI, what's not?
 - Sometimes overused and hyped buzzword
- Lack of precise definition =>
 - Difficult to design benchmarks that cannot be gamed
 - Performance on specific benchmark does not necessarily indicate general performance
 - Can lead to inflated claims or interpretations
- R. J. Sternberg [1]: “Viewed narrowly, there seem to be almost as many definitions of intelligence as there were experts asked to define it”. Legg & Hutter, “A Collection of Definitions of Intelligence”, lists over 70 definitions[2]

Roughly, two kinds of AI

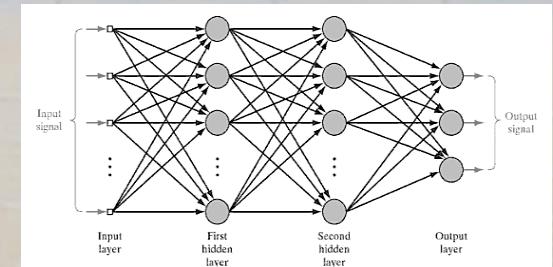
- **Symbolic (Good Old Fashion AI)**
 - **Objects are represented with symbolic data structures, behavior is explicitly programmed**
 - Objects: Attributes and methods
 - Behavior: if x in animals then ... class animals, subclass dog, cat, ...
control flow if, while, for, ...
- **Subsymbolic, neuroscience inspired**
 - **Objects are represented over a large number of simple interconnected processing units, symbols and behavior emerge**

```
dogs = [0.2 0 0 0.6 0 0 0 0 3.2 0 0 0 0]
```

```
cats = [0.3 0 0 1.2 0 0 0 0 2.1 0 0 0 4.3]
```

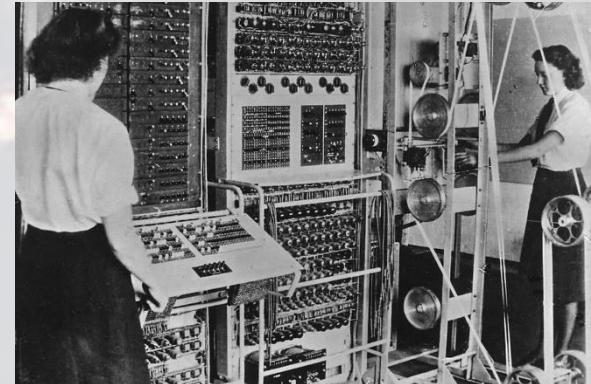
Symbols and learned behavior emerges

```
animals == [0.7 0 0 1.8 0 0 0 0 0 .3 0 0 .14 ]
```

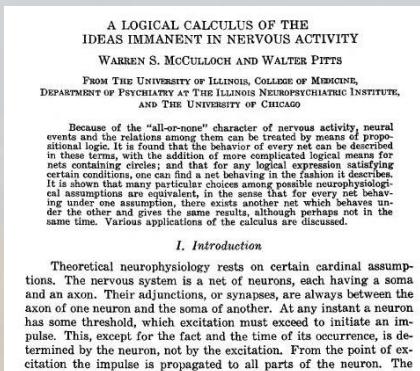


AI History

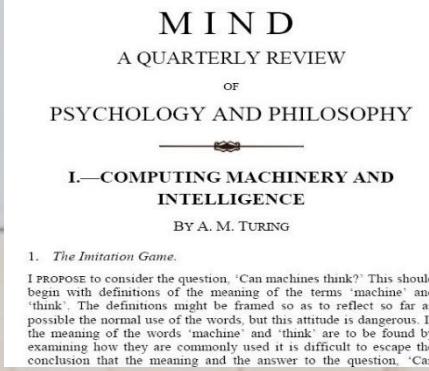
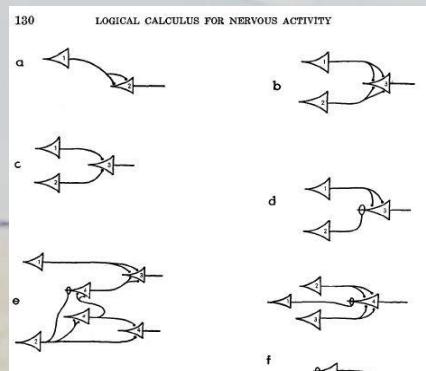
- Idea of intelligent machines almost or as old as computers
- Two approaches as old:
 - Turing machines, 1936, Turing Test, 1950 (“Symbol Processing”)
 - McCulloch and Pitts, 1943 “Subsymbolic Processing, NeuroScience inspired” :



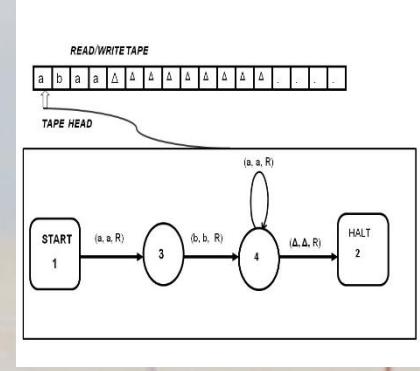
Collosus Computer, 1943



McCulloch and Pitts 1943



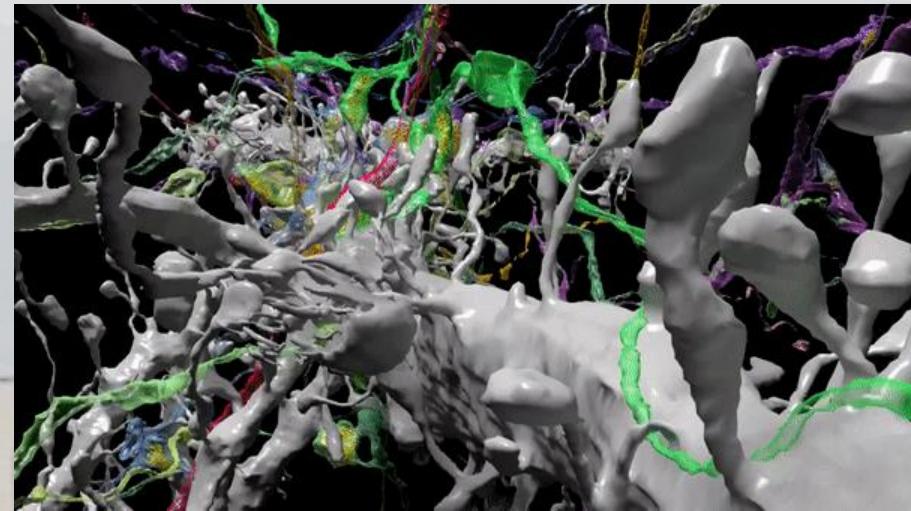
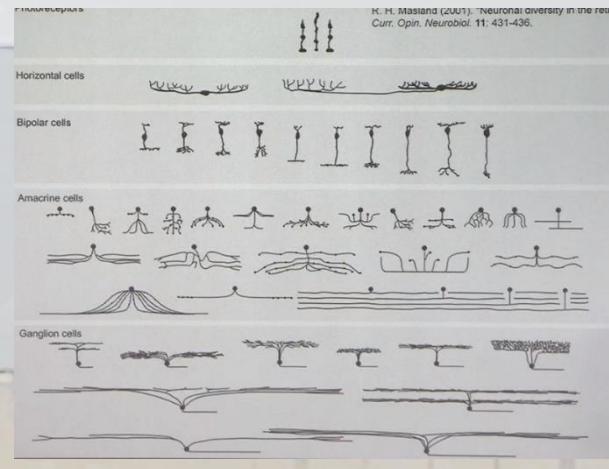
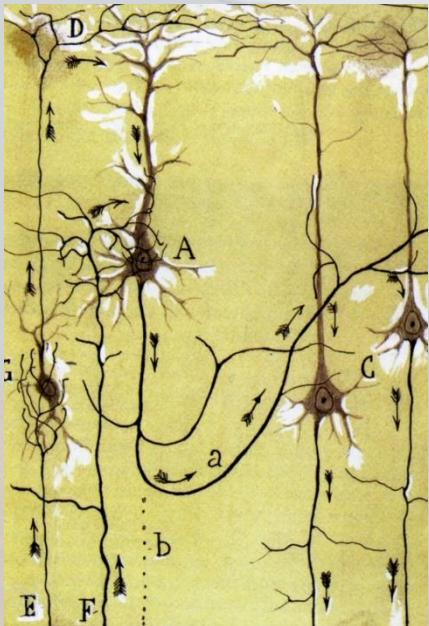
Turing, 1936



Peter Dey

SubSymbolic: Brain Inspired

- Brain: ~ 86 * billion neurons, 100 Trillion synapses
- 100,000 miles nerve fibers. And ... an absolute mess!
- Very poorly understood



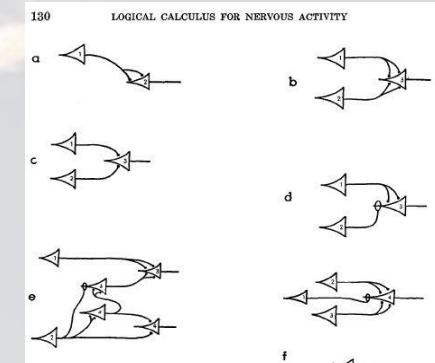
Connectomics: 1 mm³ from actual microscopy of 30nm thick slices of actual rat brain. Lichtmann

Golgi, Cajal 1906 Nobel Prize

Early Models

- McCulloch and Pitts Model, 1943

Fixed connections, neurons fire (activate) when a sufficient number of connected (via synapses) neurons fire



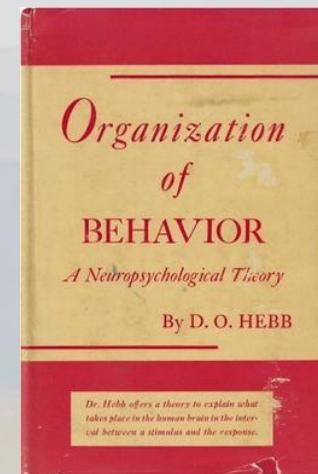
McCulloch and Pitts 1943

- Hebb, 1949: Learning "Cells that fire together wire together."

$$\Delta w_{ij} = \eta \cdot out_j \cdot in_i$$

- Modern form: Come up with an Objective function measuring output errors. (Euclidian distance of response vs desired responses for instance)

- Change weights (learning) to minimize objective function

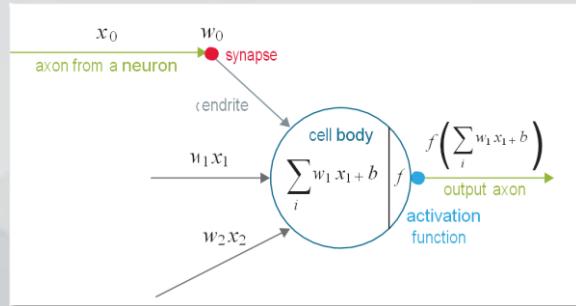


Hebb 1949

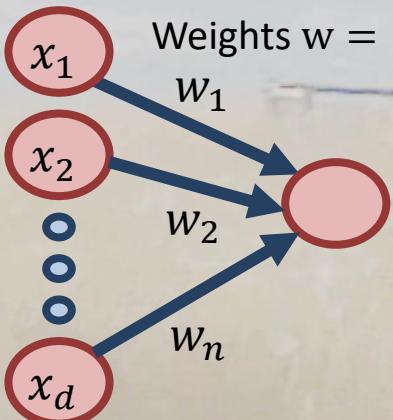
Perceptrons, 1957

- Frank Rosenblatt, Cornell
- Analog computer implementation

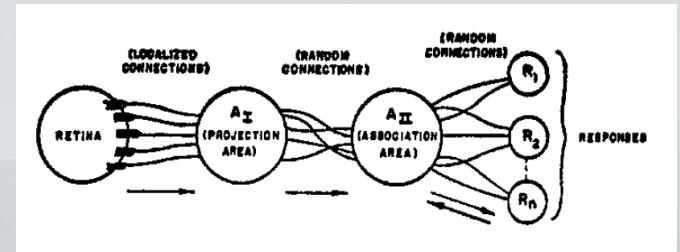
Modern formalization:



$$\text{Input } \mathbf{x} = (x_1, \dots, x_d)$$



$$\text{Output } f(\mathbf{x}) = \begin{cases} 1 & (\mathbf{x} * \mathbf{w}^T) + b > 0 \\ 0 & \text{otherwise} \end{cases}$$



Psychological Review
Vol. 65, No. 6, 1958

THE PERCEPTRON: A PROBABILISTIC MODEL FOR INFORMATION STORAGE AND ORGANIZATION IN THE BRAIN¹

F. ROSENBLATT
Cornell Aeronautical Laboratory

If we are eventually to understand the capability of higher organisms for perceptual recognition, generalization, recall, and thinking, we must first have answers to three fundamental questions:

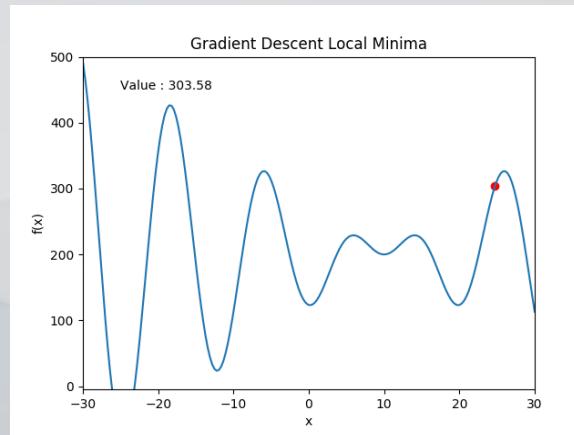
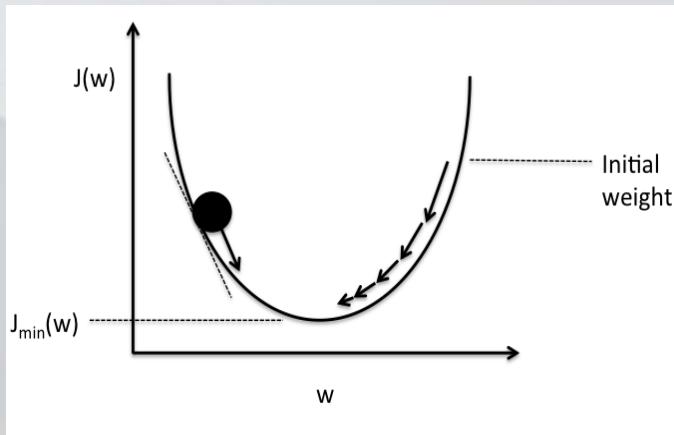
- How is information about the physical world sensed, or detected, by the biological system?
- In what form is information stored, or remembered?
- How does information contained in storage, or in memory, influence recognition and behavior?

The first of these questions is in the province of sensory physiology, and is the only one for which appreciable understanding has been achieved. This article will be concerned primarily with the second and third questions, which are still subject to a

and the stored pattern. According to this hypothesis, if one understood the code or "wiring diagram" of the nervous system, one should, in principle, be able to discover exactly what an organism remembers by reconstructing the original sensory patterns from the "memory traces" which they have left, much as we might develop a photographic negative, or translate the pattern of electrical charges in the "memory" of a digital computer.

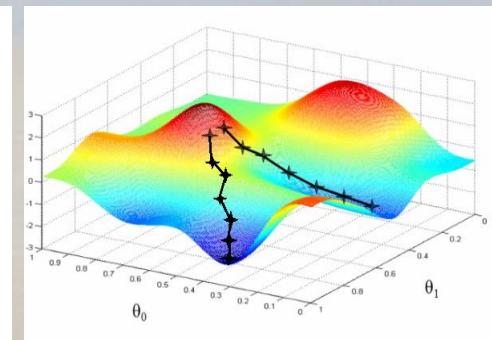
Learning implementation: Gradient Descent in weight Space

- Gradient descent: Minimize objective function $F(w)$ by subtracting from w a fraction (learning rate) of $F(w)$ derivative with respect to w (gradient)



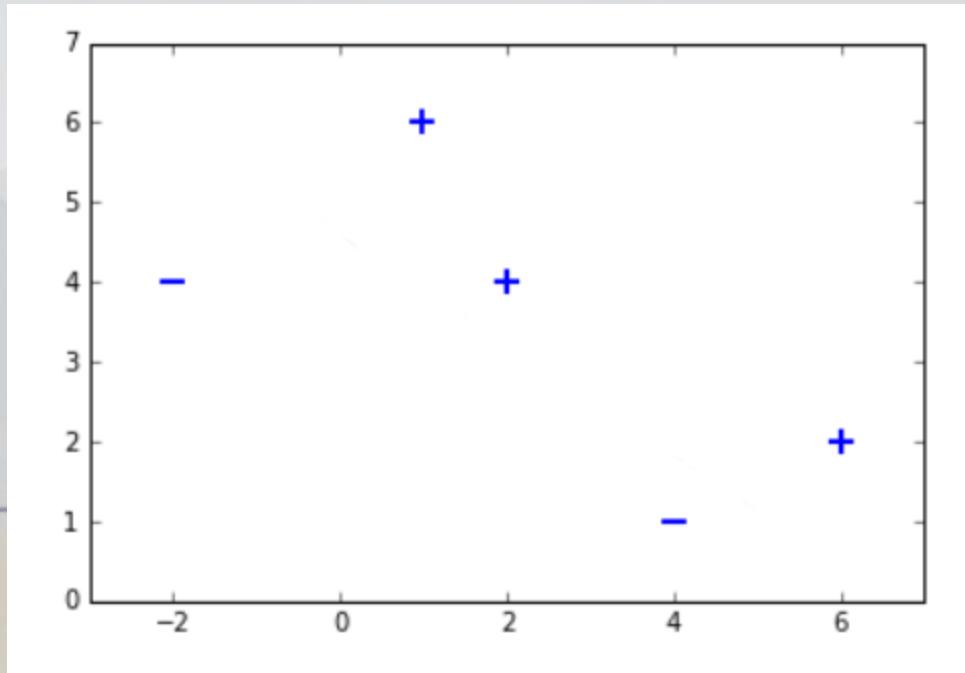
- Generalization to multi dimensional space
 $y = f(x)$, gradient of y with respect to x is the Jacobian matrix J of partial derivatives:

$$J = \begin{pmatrix} \frac{\partial y_1}{\partial x_1} & \dots & \frac{\partial y_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial y_m}{\partial x_1} & \dots & \frac{\partial y_m}{\partial x_n} \end{pmatrix}$$

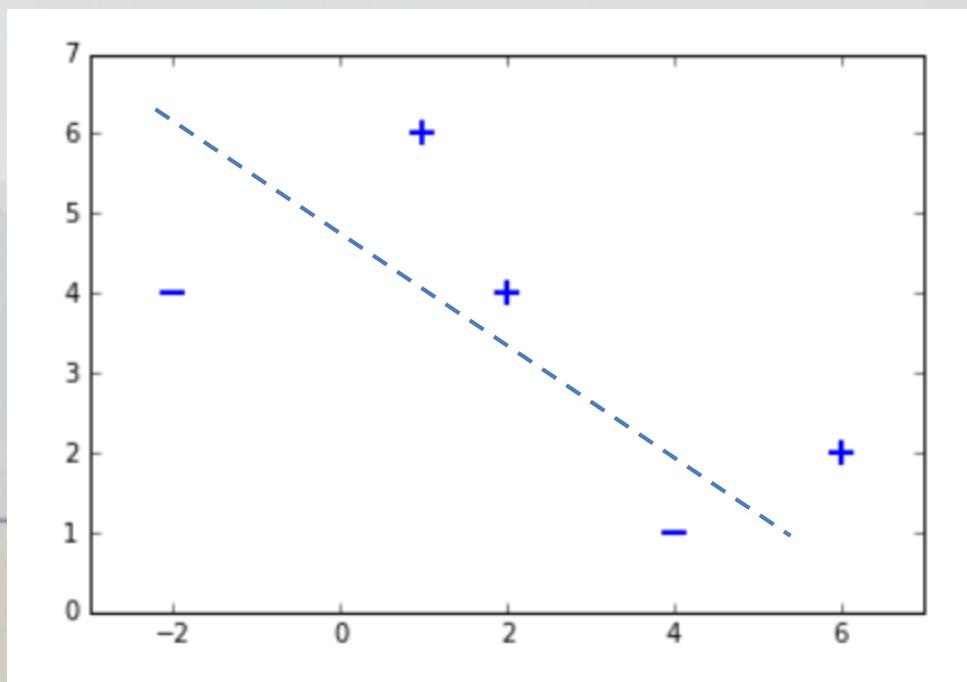
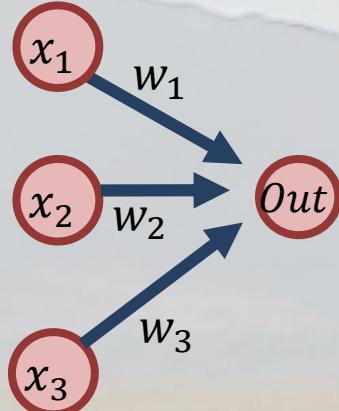


Perceptron: Python Linear Classifier Toy Example

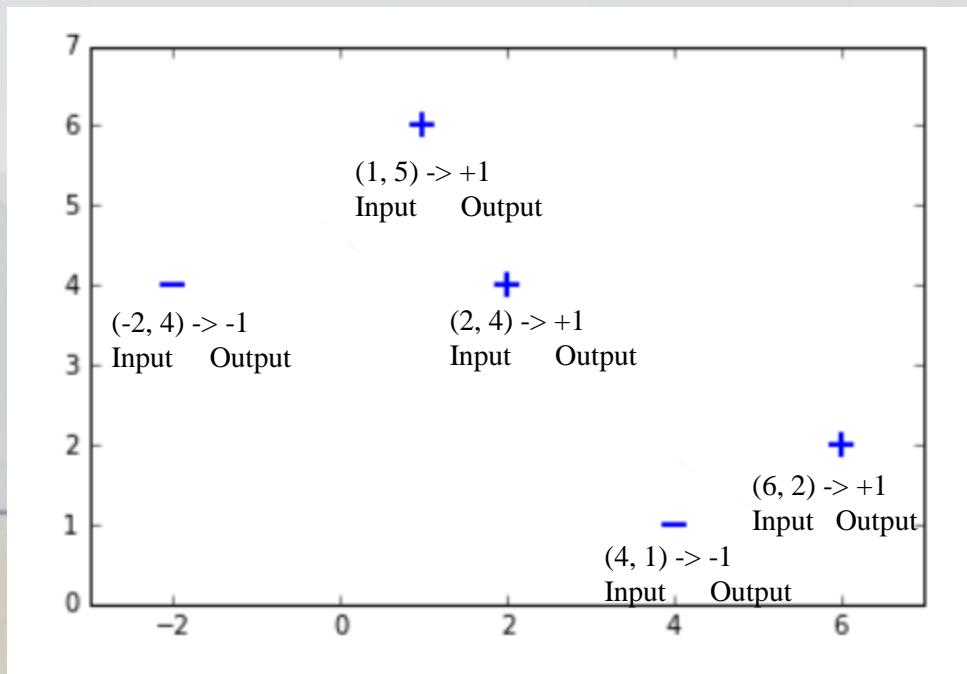
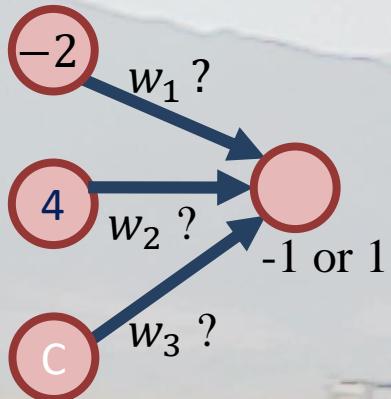
Problem: Learn to discriminate between + and - points



Perceptron: Python Linear Classifier Toy Example

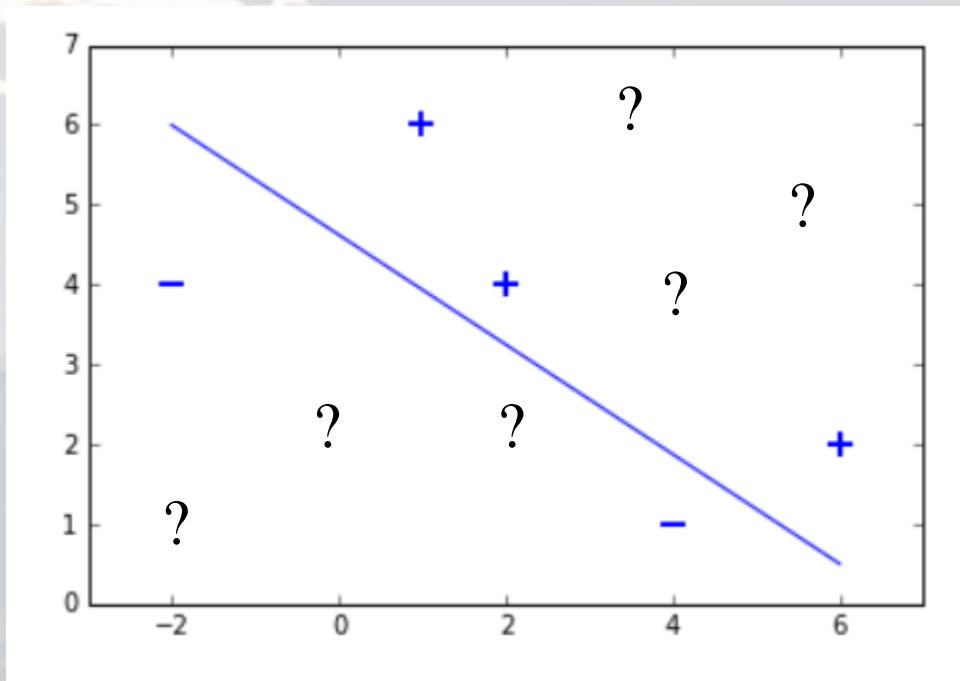


Perceptron: Python Linear Classifier Toy Example



Perceptron: Python Linear Classifier Toy Example

Test:



- Perceptron convergence Theorem (Informal description):
If there is a solution, perceptron learning will converge to the solution.

Perceptron: Python implementation with gradient descent learning

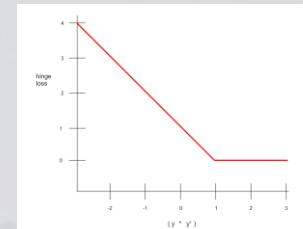
Note: We use hinge loss for objective function for historical reason

$$HL(\text{Output}) = \max(0, (1 - \text{DesiredOutput} * \text{Output}))$$

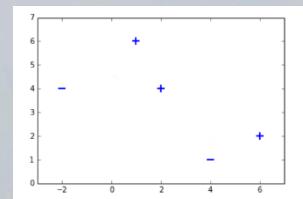
```
import numpy as np

X = np.array([ [-2,4,-1],[4,1,-1],[1,6,-1],[2,4,-1],[6,2,-1]])
# 5 inputs (points in 2d + -1 bias)
y = np.array([-1,-1,1,1,1]) # 5 Targets (Plot + or - labels)

def perceptron(X, Y):
    w = np.zeros(len(X[0])) # Weights to be modified
    eta = 1
    epochs = 20
    for e in range(epochs):
        for i, x in enumerate(X):
            if (np.dot(X[i],w)*Y[i]) <= 0:
                w = w + eta*X[i]*Y[i] #Update w with -gradient
    return w
w = perceptron(X,y)
print(w) #Weights solution
```



Loss function



After <https://github.com/MaviccPRP/perceptron>

Perceptron: Python implementation with gradient descent learning

```
import numpy as np

X = np.array([ [-2,4,-1],[4,1,-1],[1,6,-1],[2,4,-1],[6,2,-1]])
# 5 inputs (points in 2d + -1 bias)
y = np.array([-1,-1,1,1,1]) # 5 Targets (Plot + or - labels)

def perceptron(X, Y):
    w = np.zeros(len(X[0])) # Weights to be modified
    eta = 1
    epochs = 20
    for e in range(epochs):
        for i, x in enumerate(X):
            if (np.dot(X[i], w) * Y[i]) <= 0:
                w = w + eta*x*Y[i]
    return w
w = perceptron(X,y)
print(w) #Weights solution
```

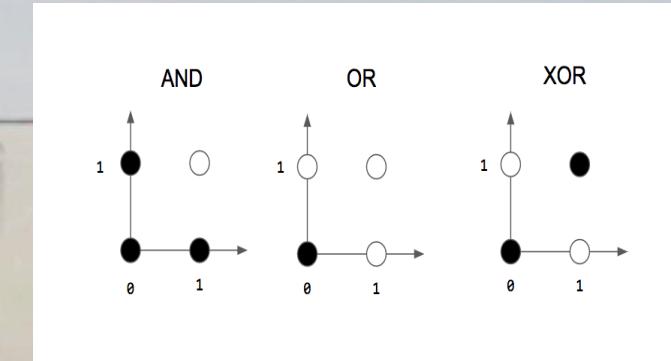
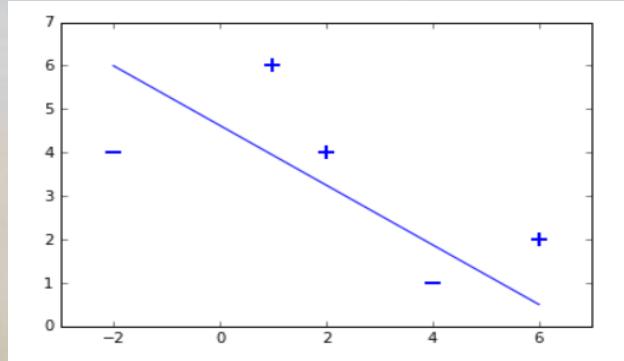
Error measured with hinge loss function:
 $HL(\text{Output}) = \max [0, (1 - \text{DesiredOutput} * \text{Output})]$

$$HL(wX[i]) = \max [0, (1 - Y[i] * wX[i])]$$

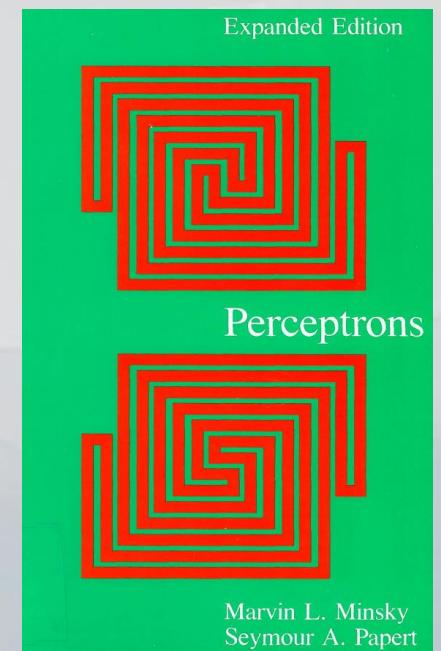
$$\begin{aligned} dHL(w)/dw &= d(\max[0, 1 - (Y[i] * wX[i])])/dw \\ &= 0 && \text{if } 1 - Y[i] * wX[i] > 0 \\ &= -\eta * (-X[i]Y[i]) && \text{if } 1 - Y[i] * wX[i] \leq 0 \end{aligned}$$

1970's Temporary end of Neural Networks!

- What about other functions?
- How about Boolean operations?
- AND ? ✓ OR ? ✓
- XOR ? ✗ No can do!
- Need linear separability
- Expectations were inflated



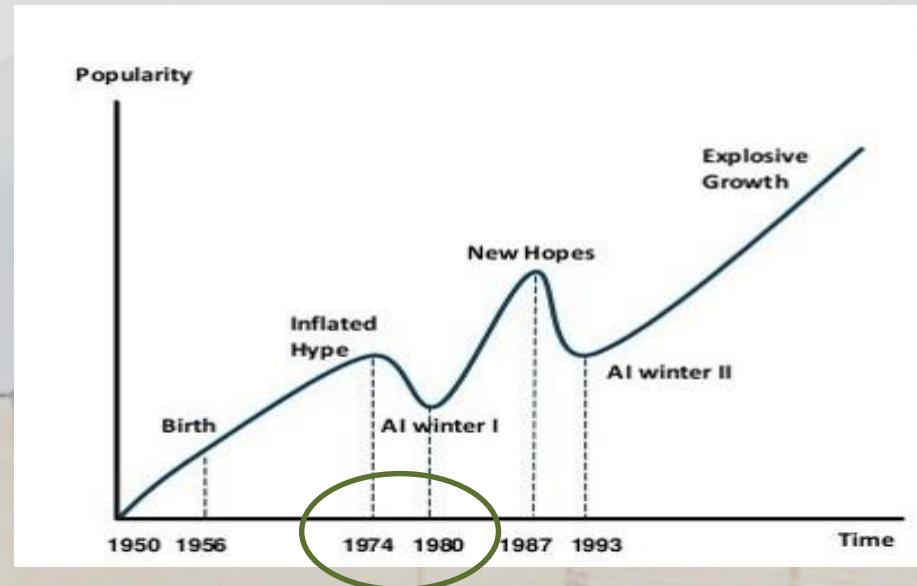
Minsky & Papert, 1969



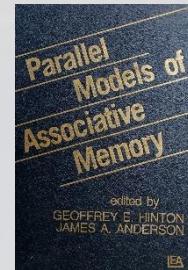
Meanwhile ... 50's onward, symbolic AI

- Development of what is now called GOFAI (Good Old Fashion AI)
- Reinforcement Learning Expert systems, probabilistic models, use of Markov decision processes, models of reasoning, natural language processing, Search algorithms, etc ...

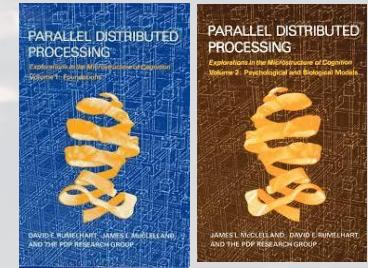
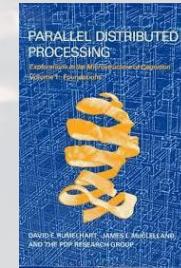
- But: Inflated expectations also!
- => First “AI Winter”:



Mid 80's, rebirth: Connectionism



Hinton and Anderson,
1981



Rumelhart, McClelland
and the PDP Group, 1986

- PDP Group: Multi disciplinary group of researchers.
- Key contributions:
- Learning with Backpropagation*, for multilayer networks, with eg sigmoid squash ,
- Recurrent network architecture allowing processing in time (Backpropagation through time)
- Restricted Boltzmann machines, Boltzmann machines, ...
- Many applications: Linguistic morphology, sentence processing, speech perception.
- Formulation of Subsymbolic (vs symbolic AI) theory of mind:

Massively distributed representations

Processed in parallel

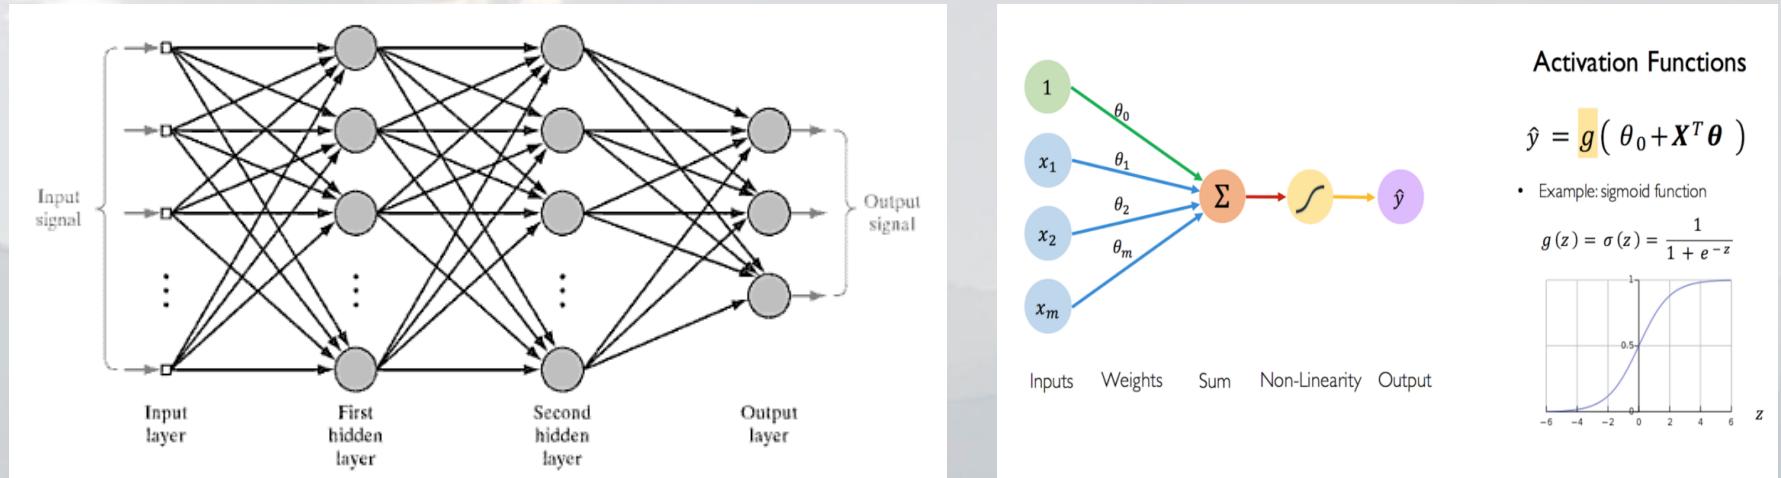
Graceful degradation

Symbols emerge from patterns of activation

Noise resistance

Multilayer Network: How can they learn?

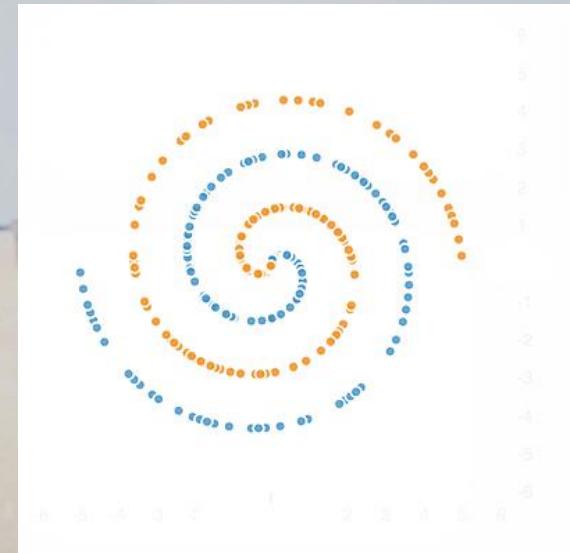
Backpropagation



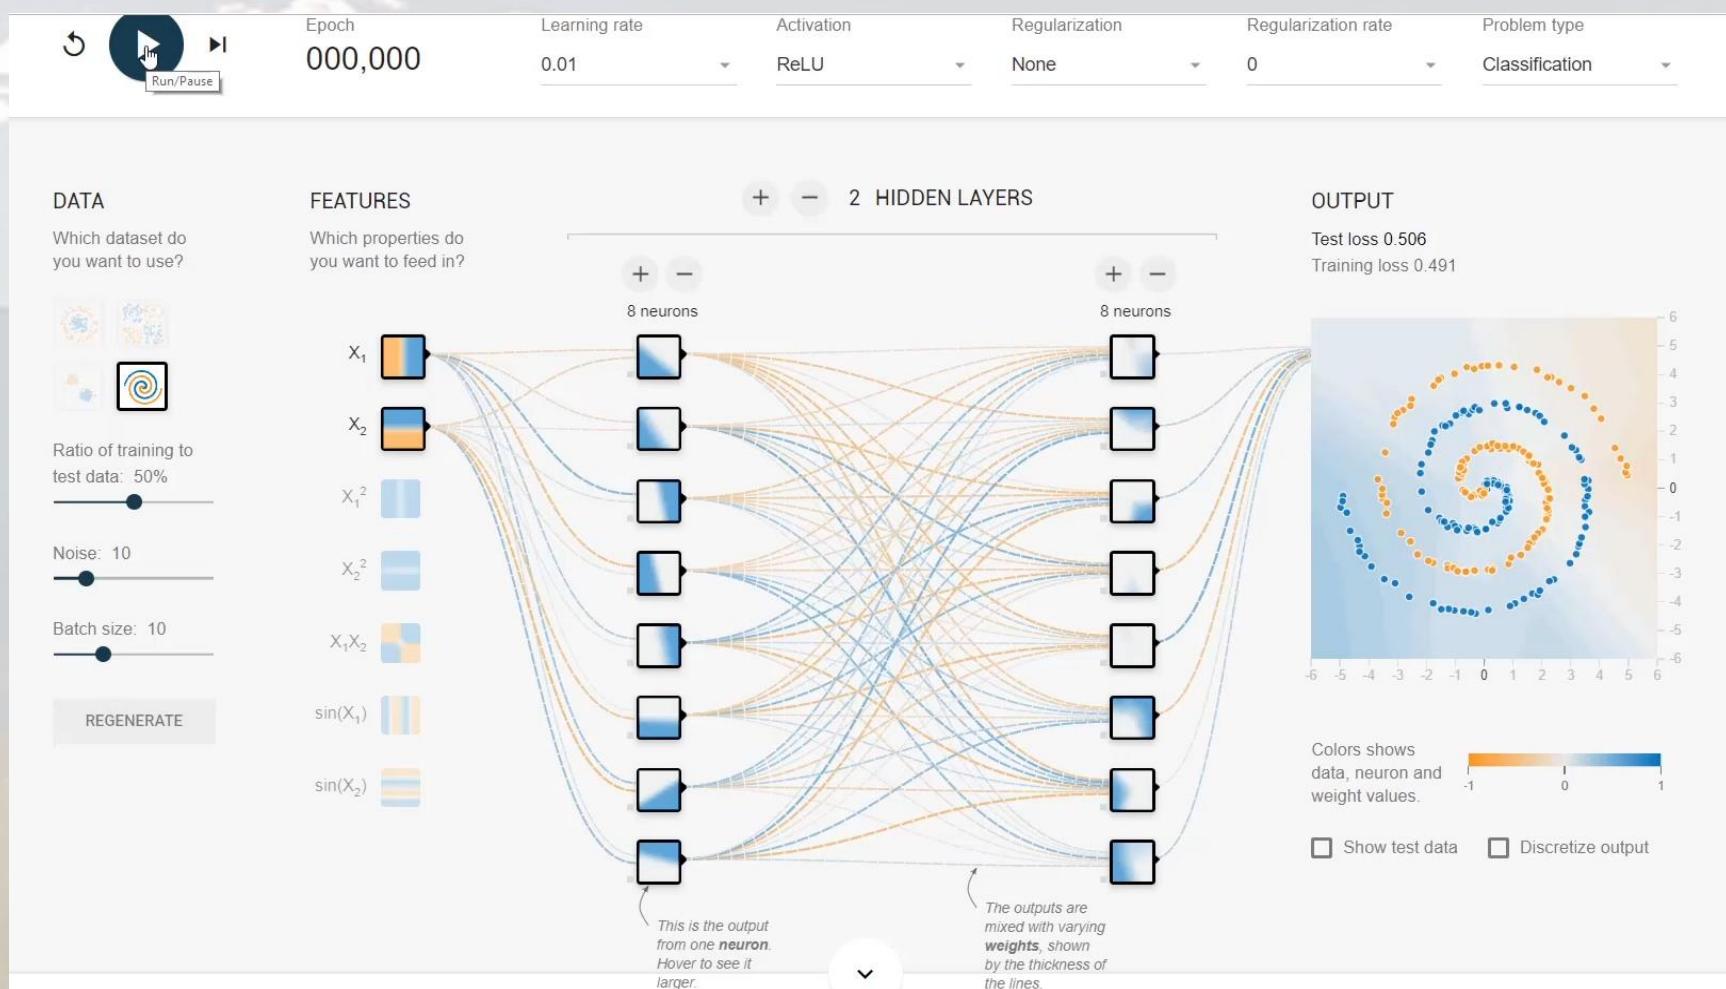
- Use sigmoid or similar to allow for gradients
- General Idea: Use the chain rule for derivatives:
- $F(x) = f(g(x)) \Rightarrow F'(x) = f'(g(x)) * g'(x)$ Use multivariable calculus: Jacobian to correct weight error.
- Store weight values on forward pass, then propagate error backward towards previous layers and change weights

Multilayer Networks

- Theoretical result:
Any continuous function can be learned with arbitrary precision with a neural network with one hidden layer (given enough hidden units)
- Harder classification Problem that the perceptron we saw earlier cannot solve:



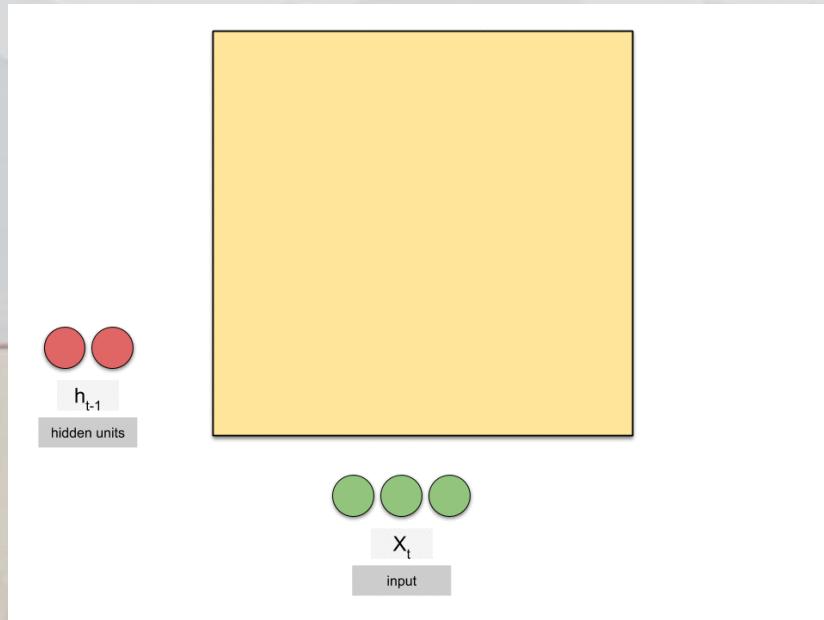
Multilayer Network: <https://playground.tensorflow.org>



What about time? Recurrent Neural Networks

- Temporal learning and processing of data sequences in time

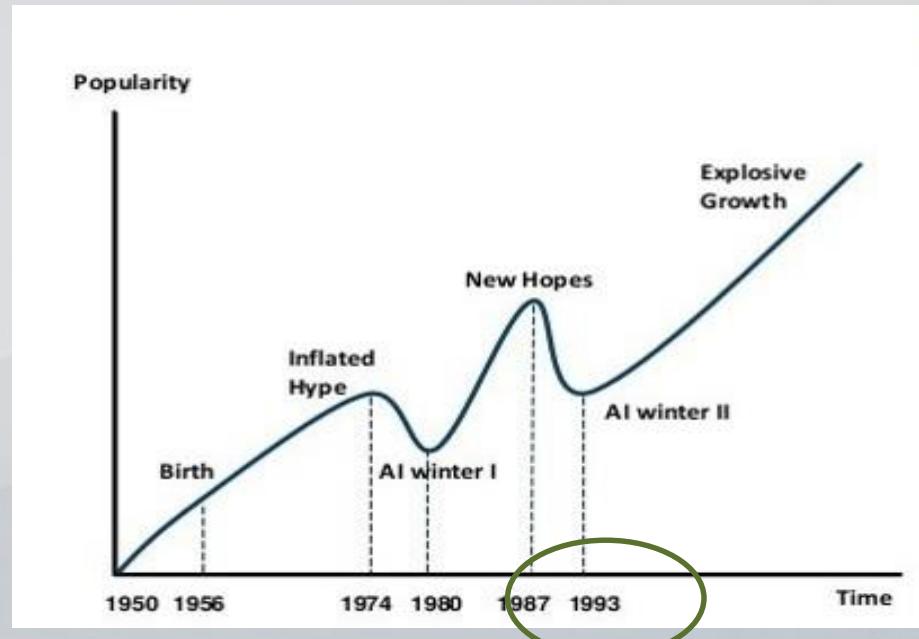
General Idea: Use recurrent connections from hidden layer back to itself; Combine input layer at time t with hidden layer at time $t-1$



Raimi Karim

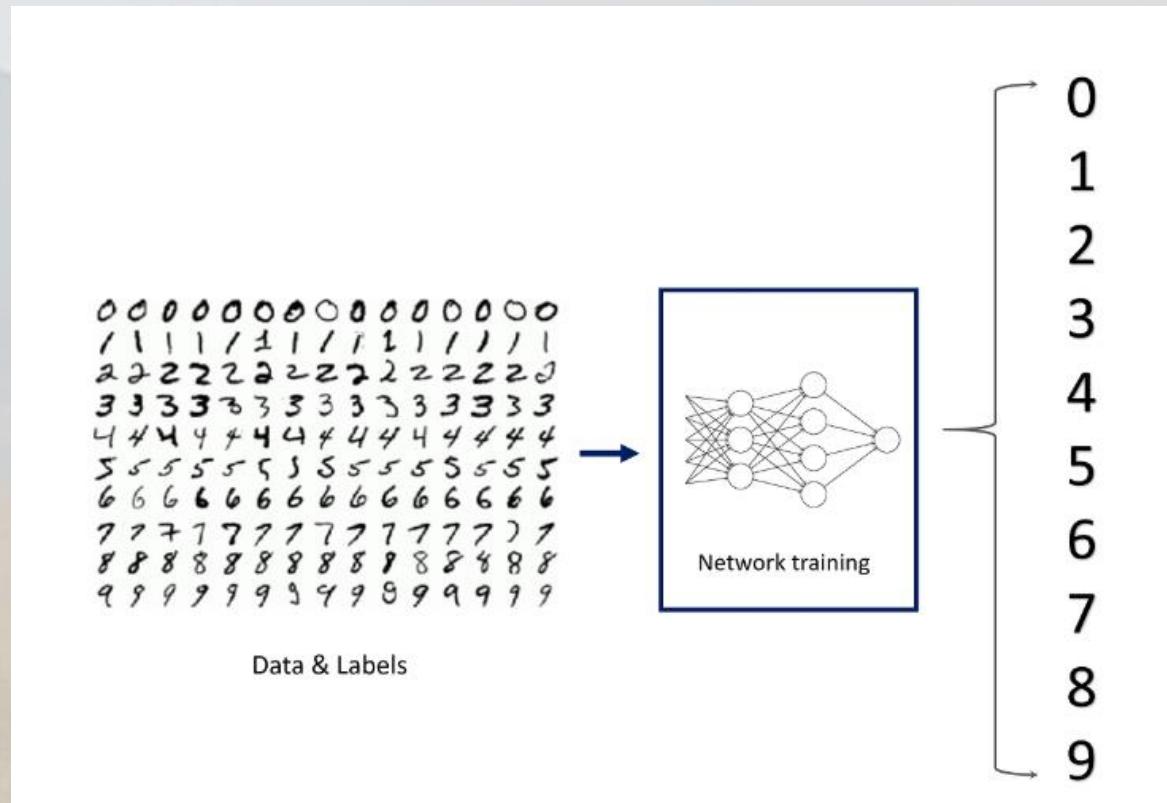
But ... Mid-90s: Second AI Winter!

- NNs are data hungry, too little data
- NNs are compute hungry, too little compute
- Hype and over promises



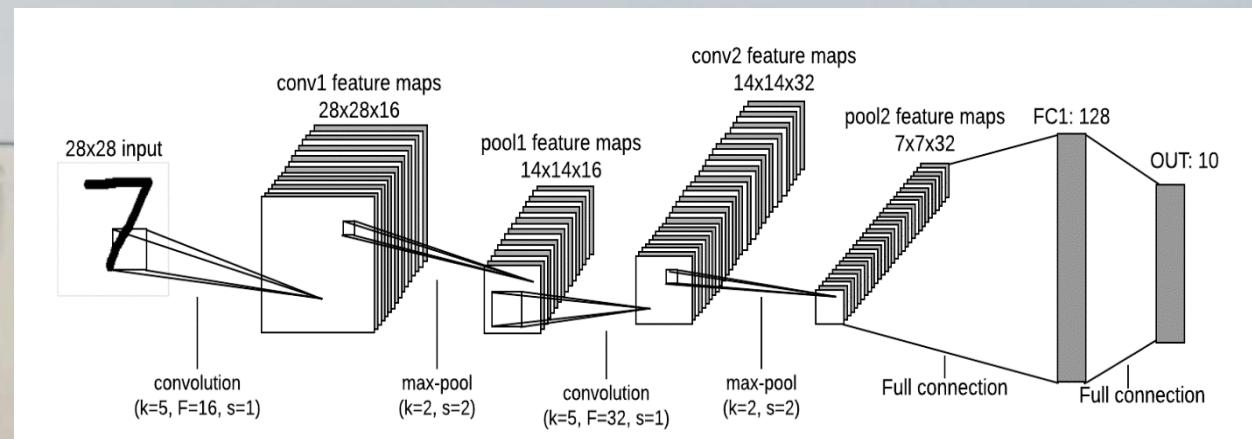
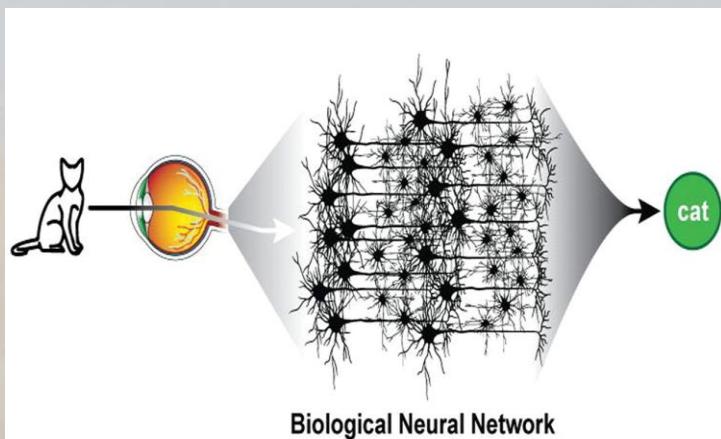
Meanwhile, however ...

- Yann Lecun, Bell Labs
- Handwritten digit recognition

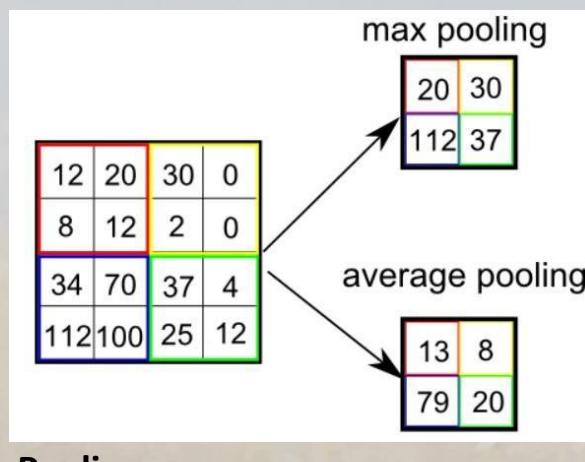
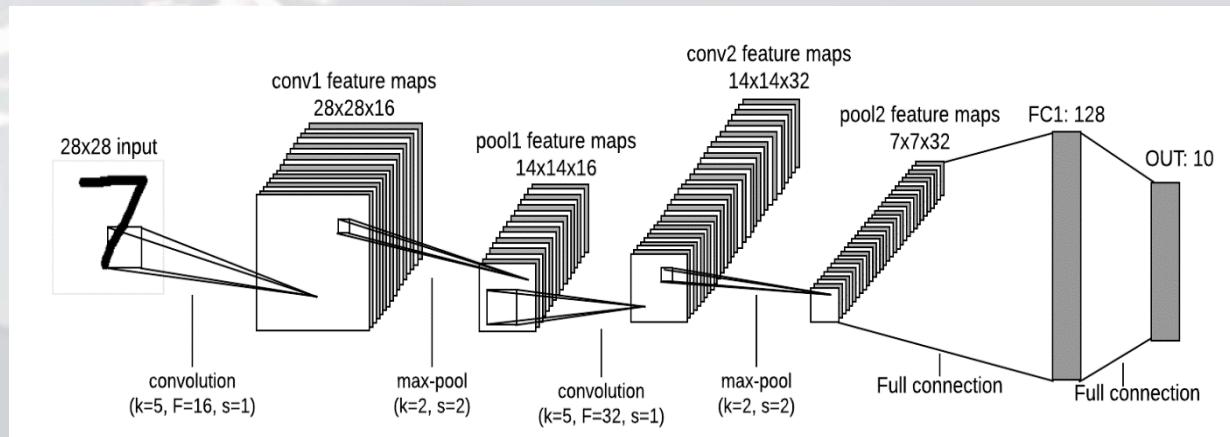


Convolution Neural Network

- Hierarchical processing, inspired by some cells in visual cortex: Some cells respond specifically to certain features like edges => Filters
- Idea: Intermediate layers detect features increasingly more complex
- Two new kind of layers: Convolution and pooling layers



Convolution Neural Network: Pooling

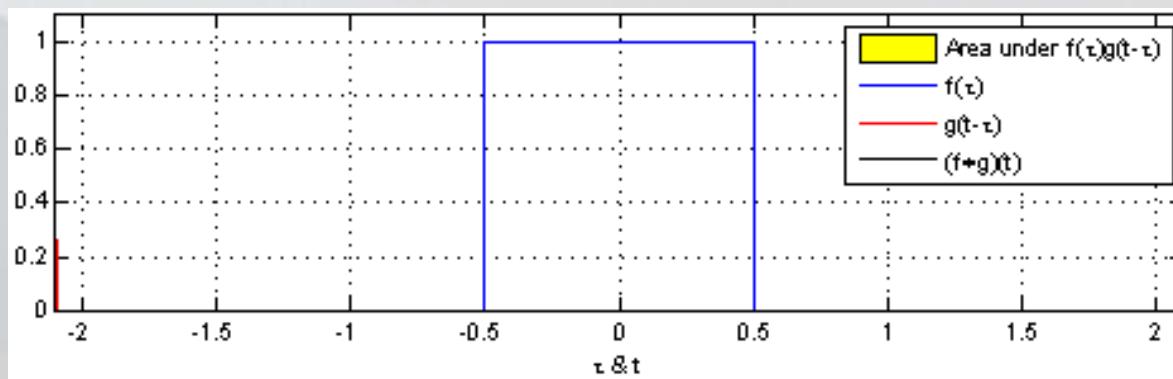


Pooling

Convolution

Central to signal processing, filters

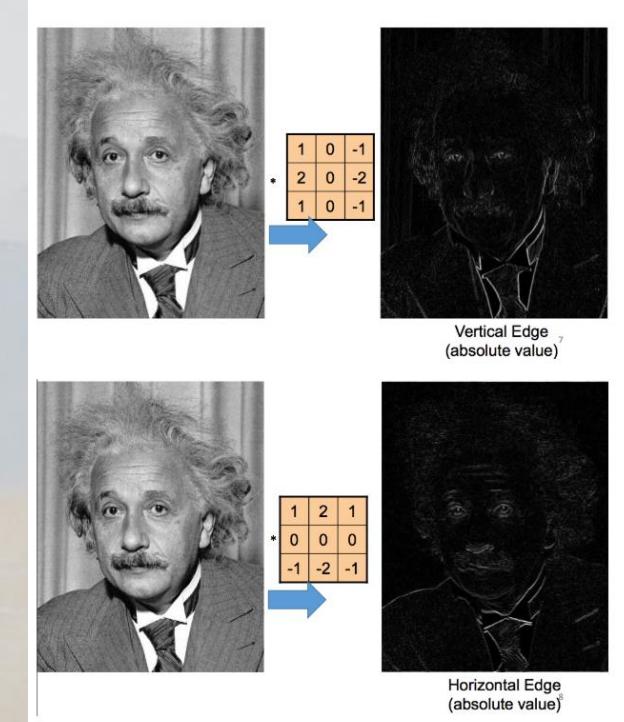
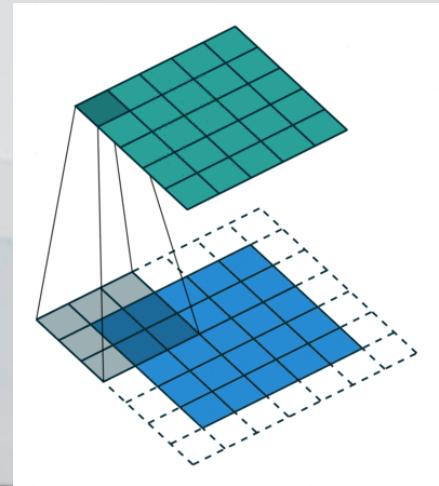
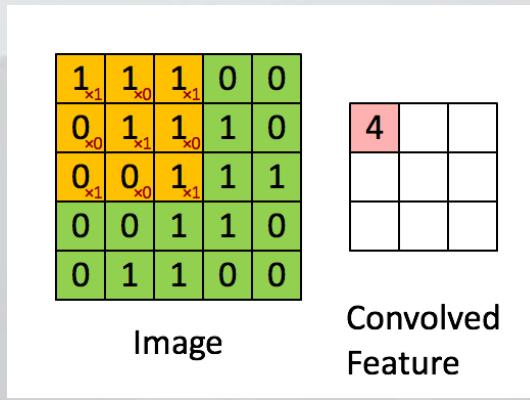
Continuous: $(f * g)(t) = \int f(\tau)g(t - \tau) d\tau$



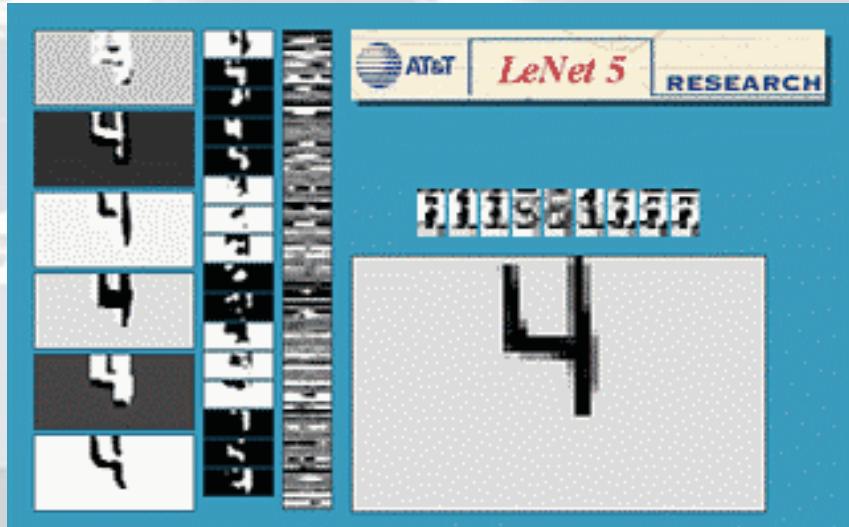
Convolutions

Continuous: $(f * g)(t) = \int f(\tau)g(t - \tau) d\tau$

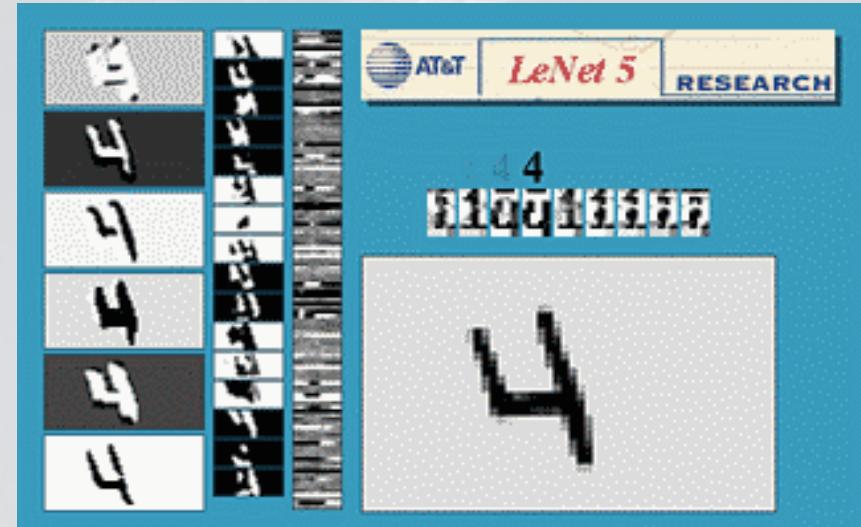
Discrete: $(f * g)[n] = \sum f[m] g[n - m]$



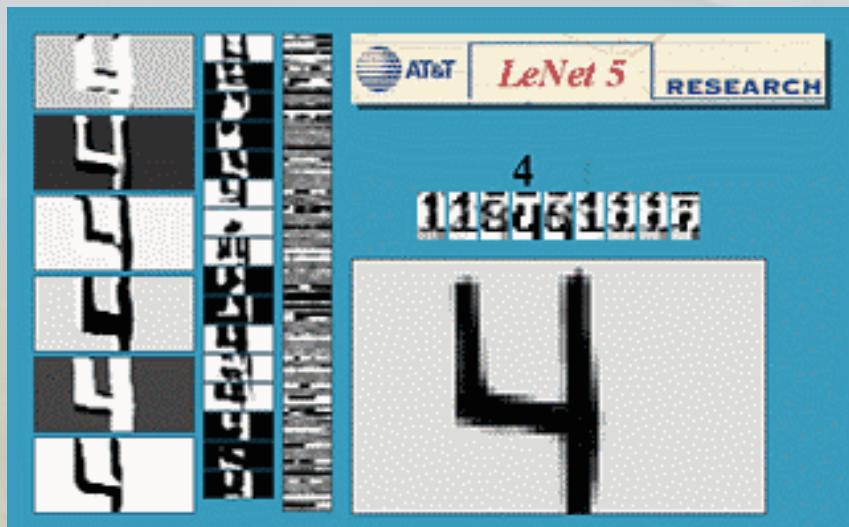
Vertical and horizontal edge filters



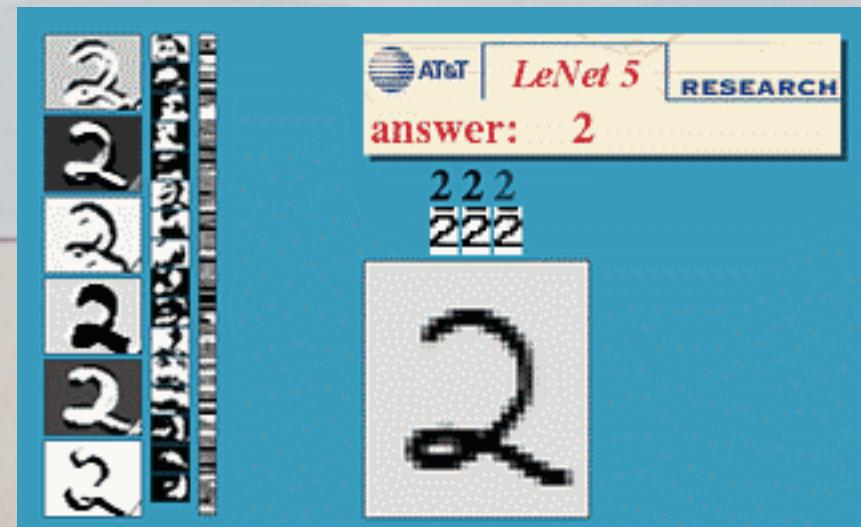
Translation Invariance



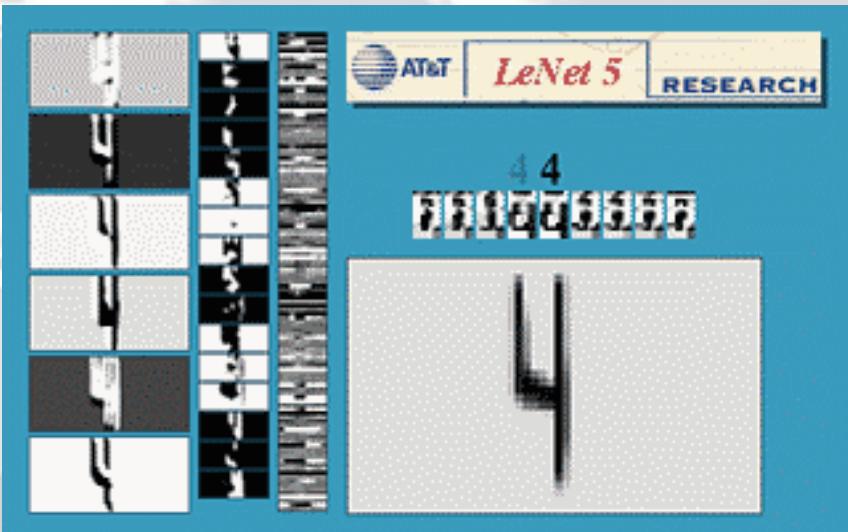
Rotation Invariance



Scale Invariance



Resistance to Noise



Aspect Ratio Invariance



Stroke Width Invariance



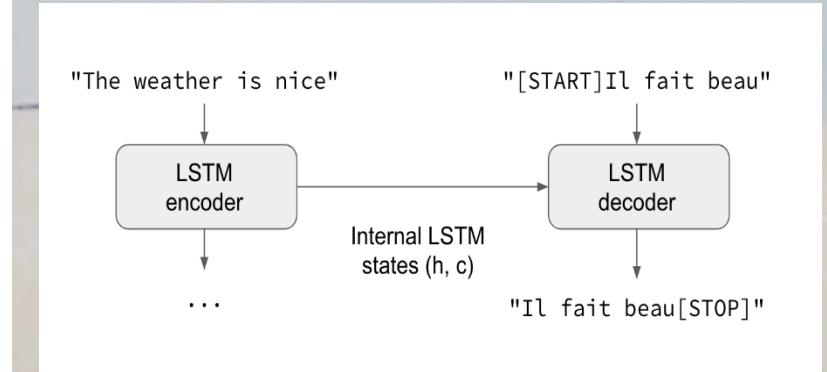
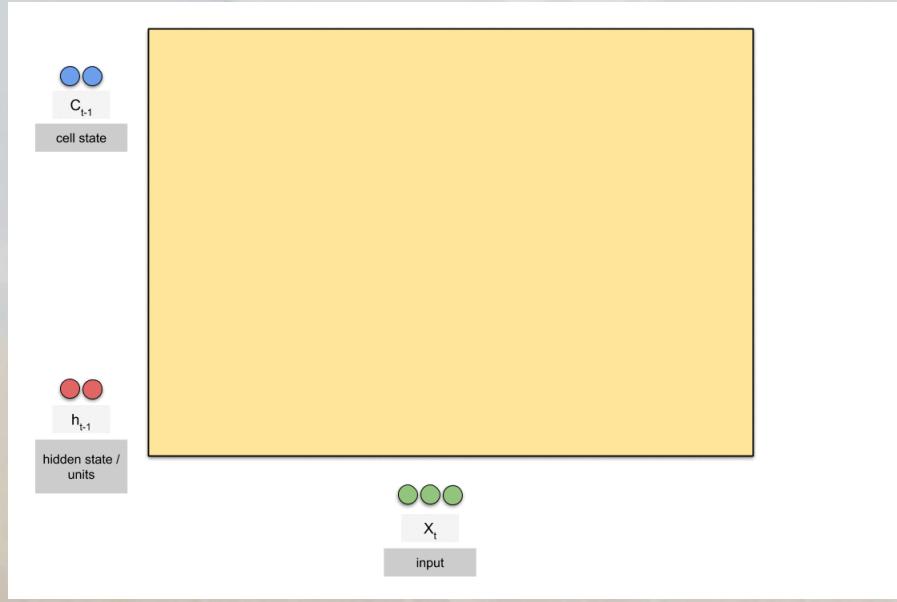
Weirdos



Weirdos

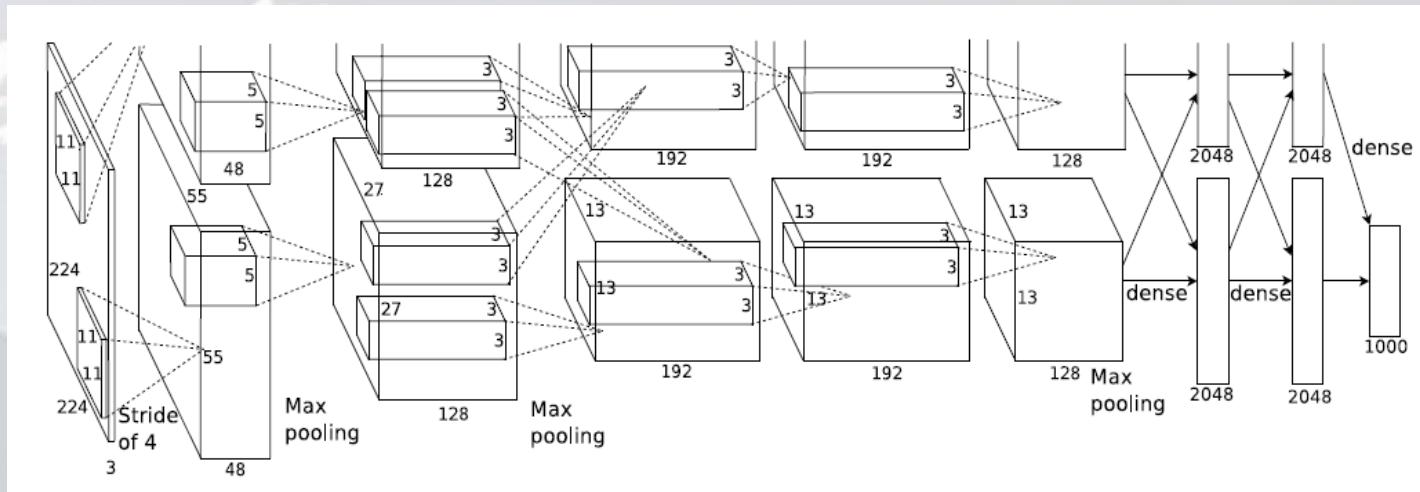
LSTMs: Solving the problem of vanishing gradients

- Hochreiter & Schmidhuber, 1991, 1995
- Long Short Term Memory. “Selectively remember or forget”
- Later used with sequence to sequence networks:

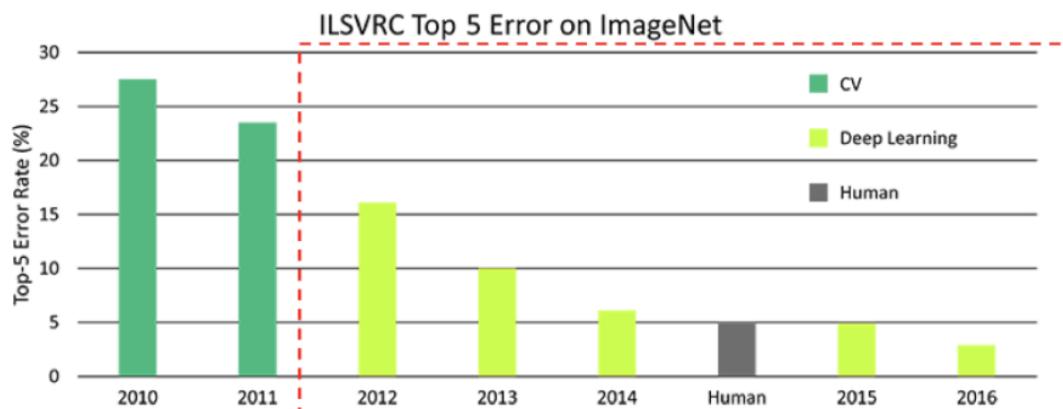


Animation credit: Raimi Karim

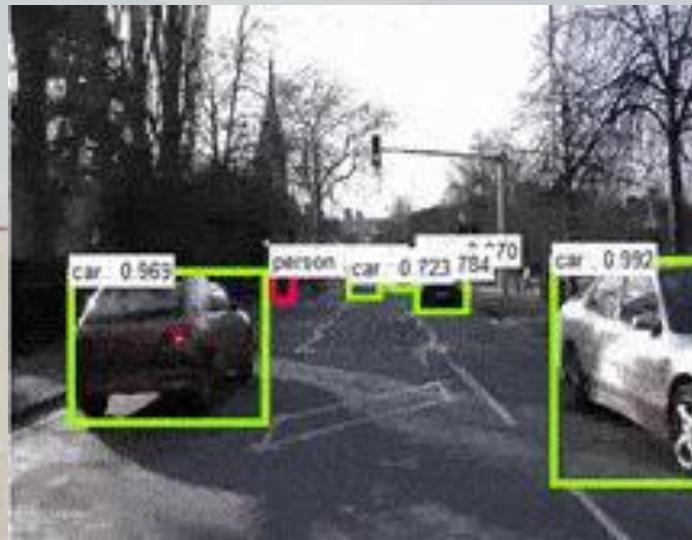
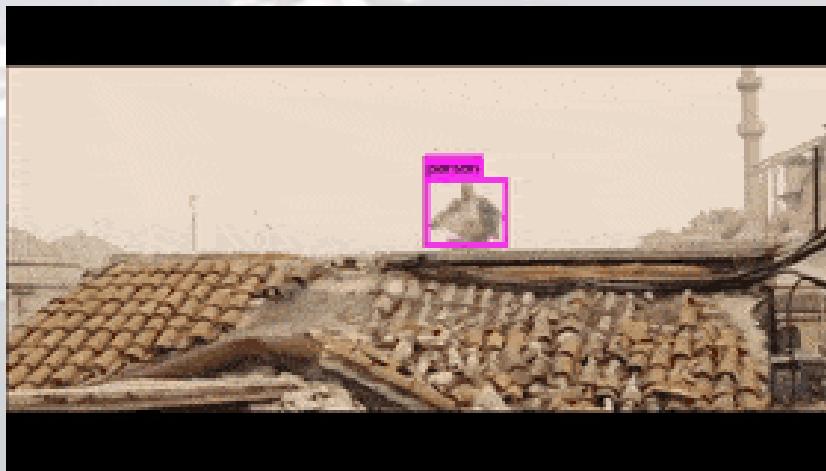
CNNs: 2012 AlexNet



CNNs: 2012 AlexNet



From stills to videos: Sliding windows

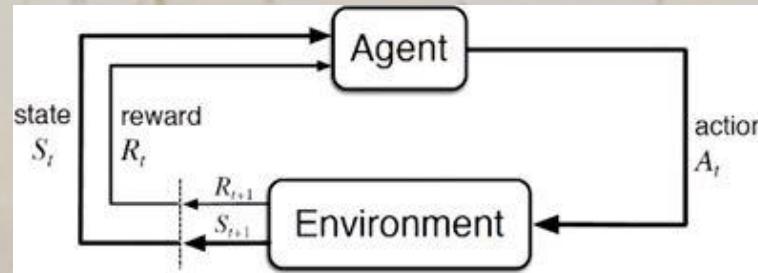


Neural network Machine Learning

- Most is supervised Learning, but let's not forget:
- Unsupervised Learning eg autoencoders
- Reinforcement Learning

Reinforcement Learning

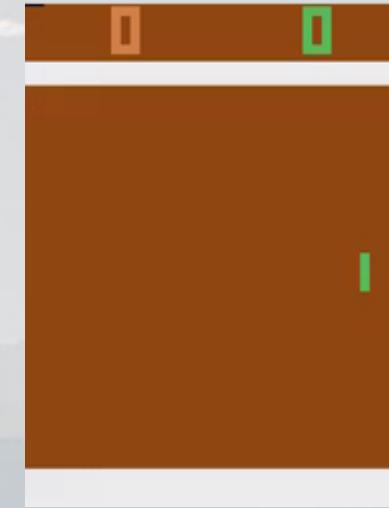
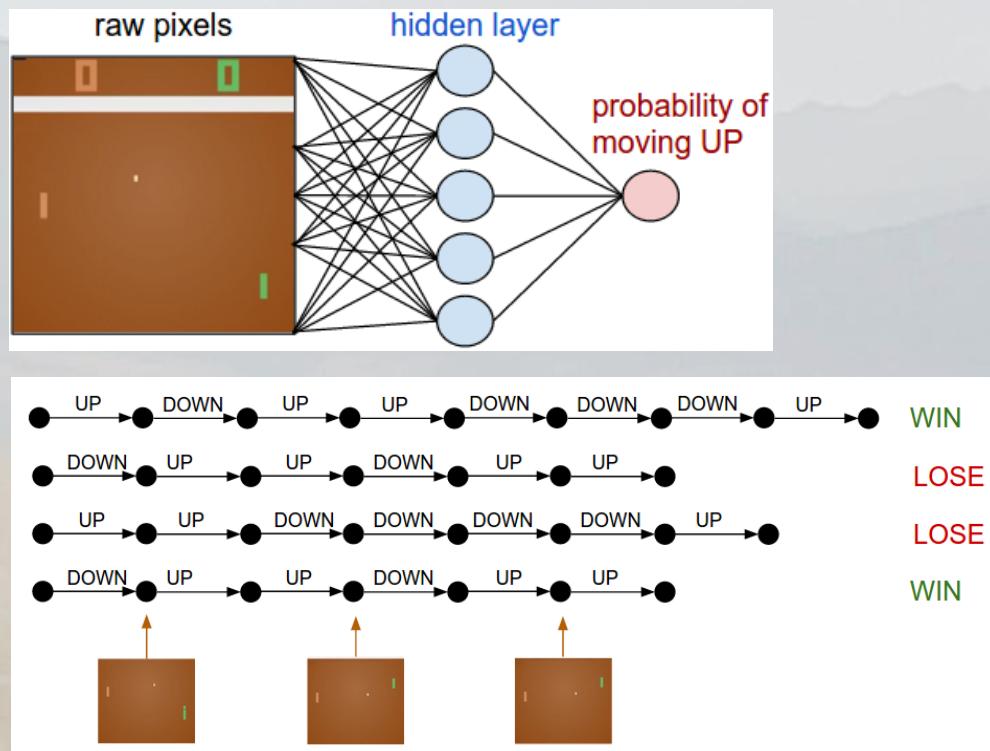
- Rich AI history from the 50's
- In its simplest form:
- Agent evolves in environment of states, performing actions leading to successive states, following a policy with rewards and penalties
- Learning consists in finding best policy maximizing rewards
- It is possible to learn policy with neural network.



Reinforcement Learning: A simple example (2016)

<http://karpathy.github.io> (Highly recommended if new to RL. Policy gradients)

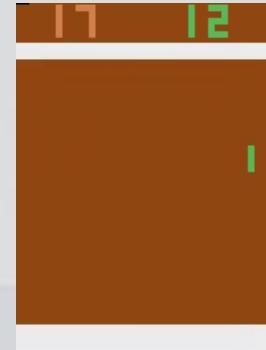
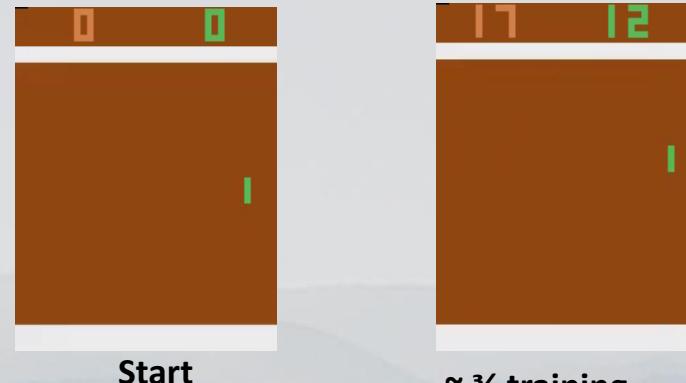
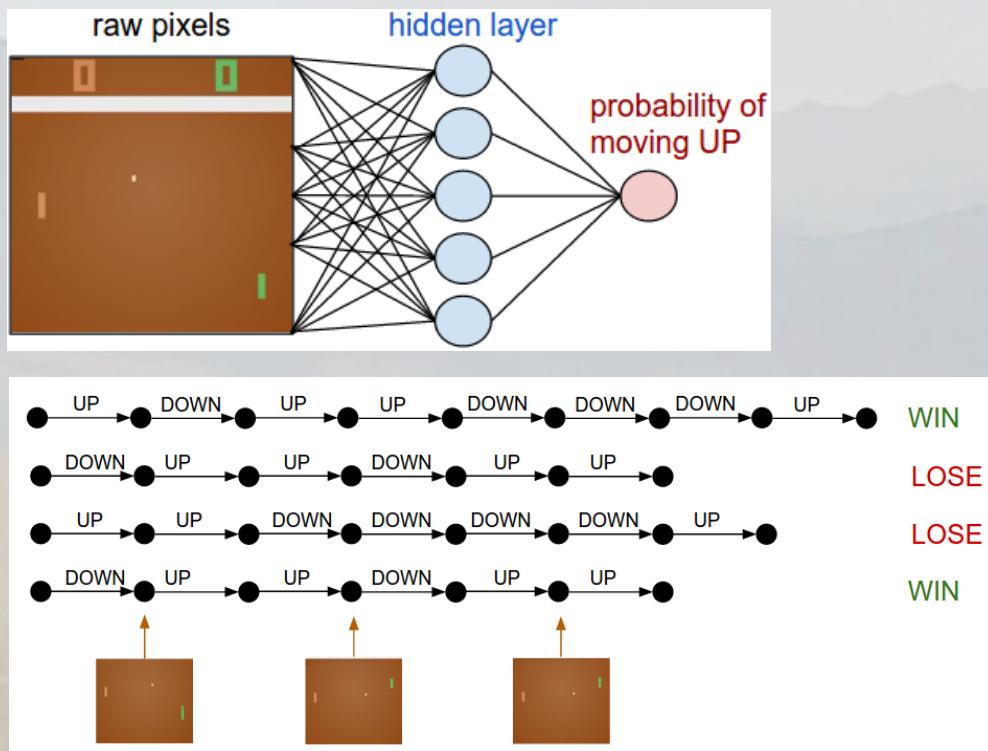
- Learn policy from raw pixels
- Well worth a close look:
130 lines Python implementation



Reinforcement Learning: A simple example (2016)

<http://karpathy.github.io> (Highly recommended if new to RL with policy gradients)

- Learn policy from raw pixels
- Well worth a close look:
130 lines Python implementation



Neural Nets: Many architectures

Residual Networks (ResNet)

GANS: Generative Adversarial Networks (two networks)

Attention networks

Auto encoders (Unsupervised learning)

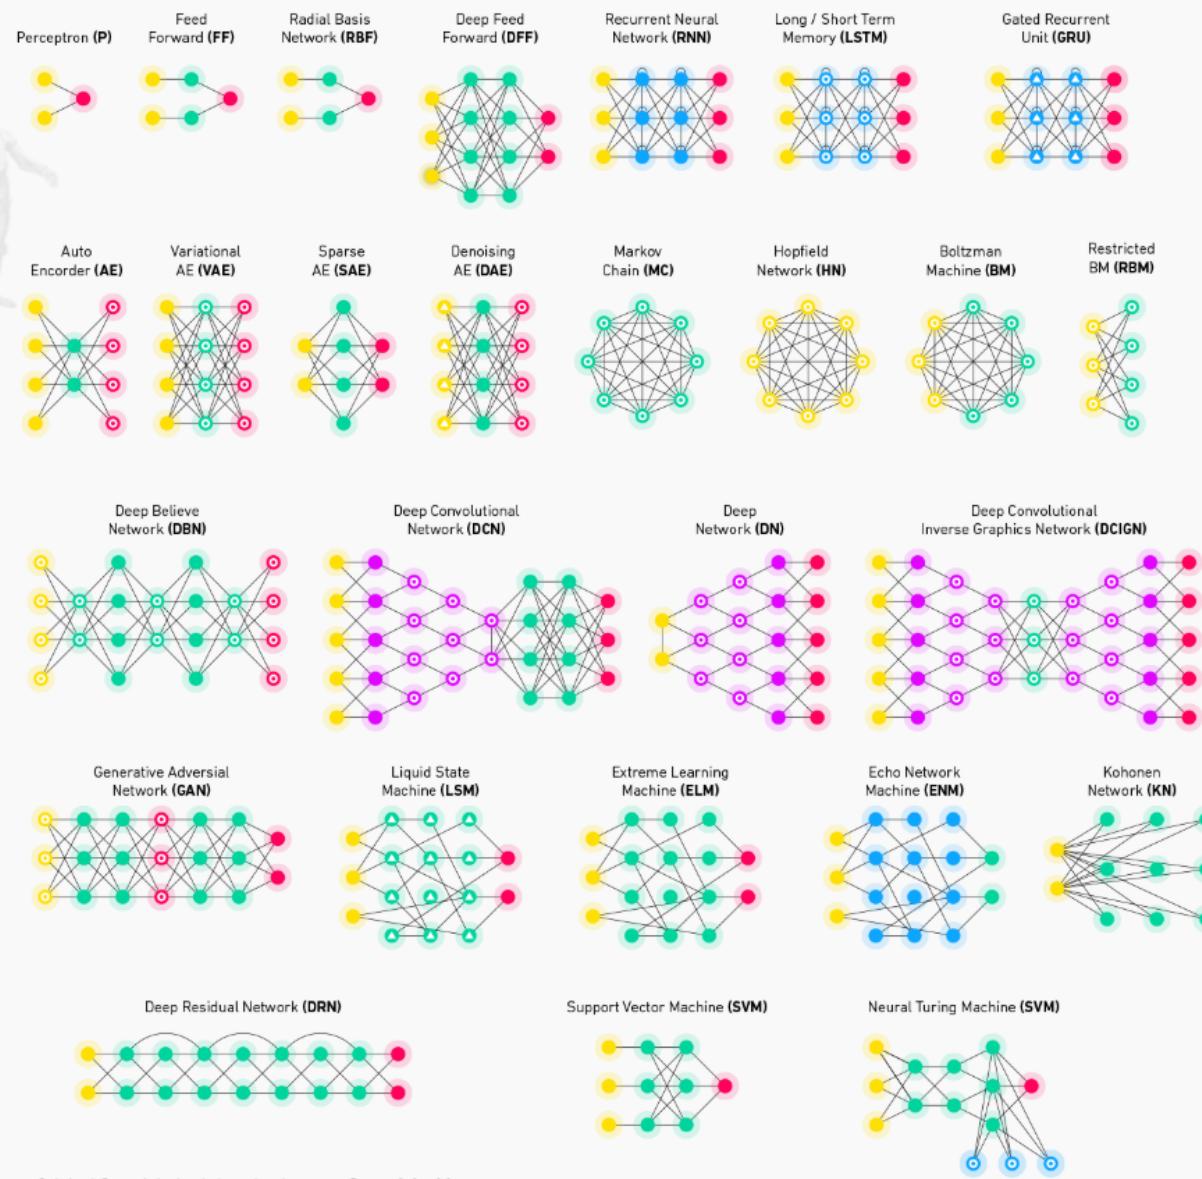
Variational Autoencoders

Restricted Boltzmann machines

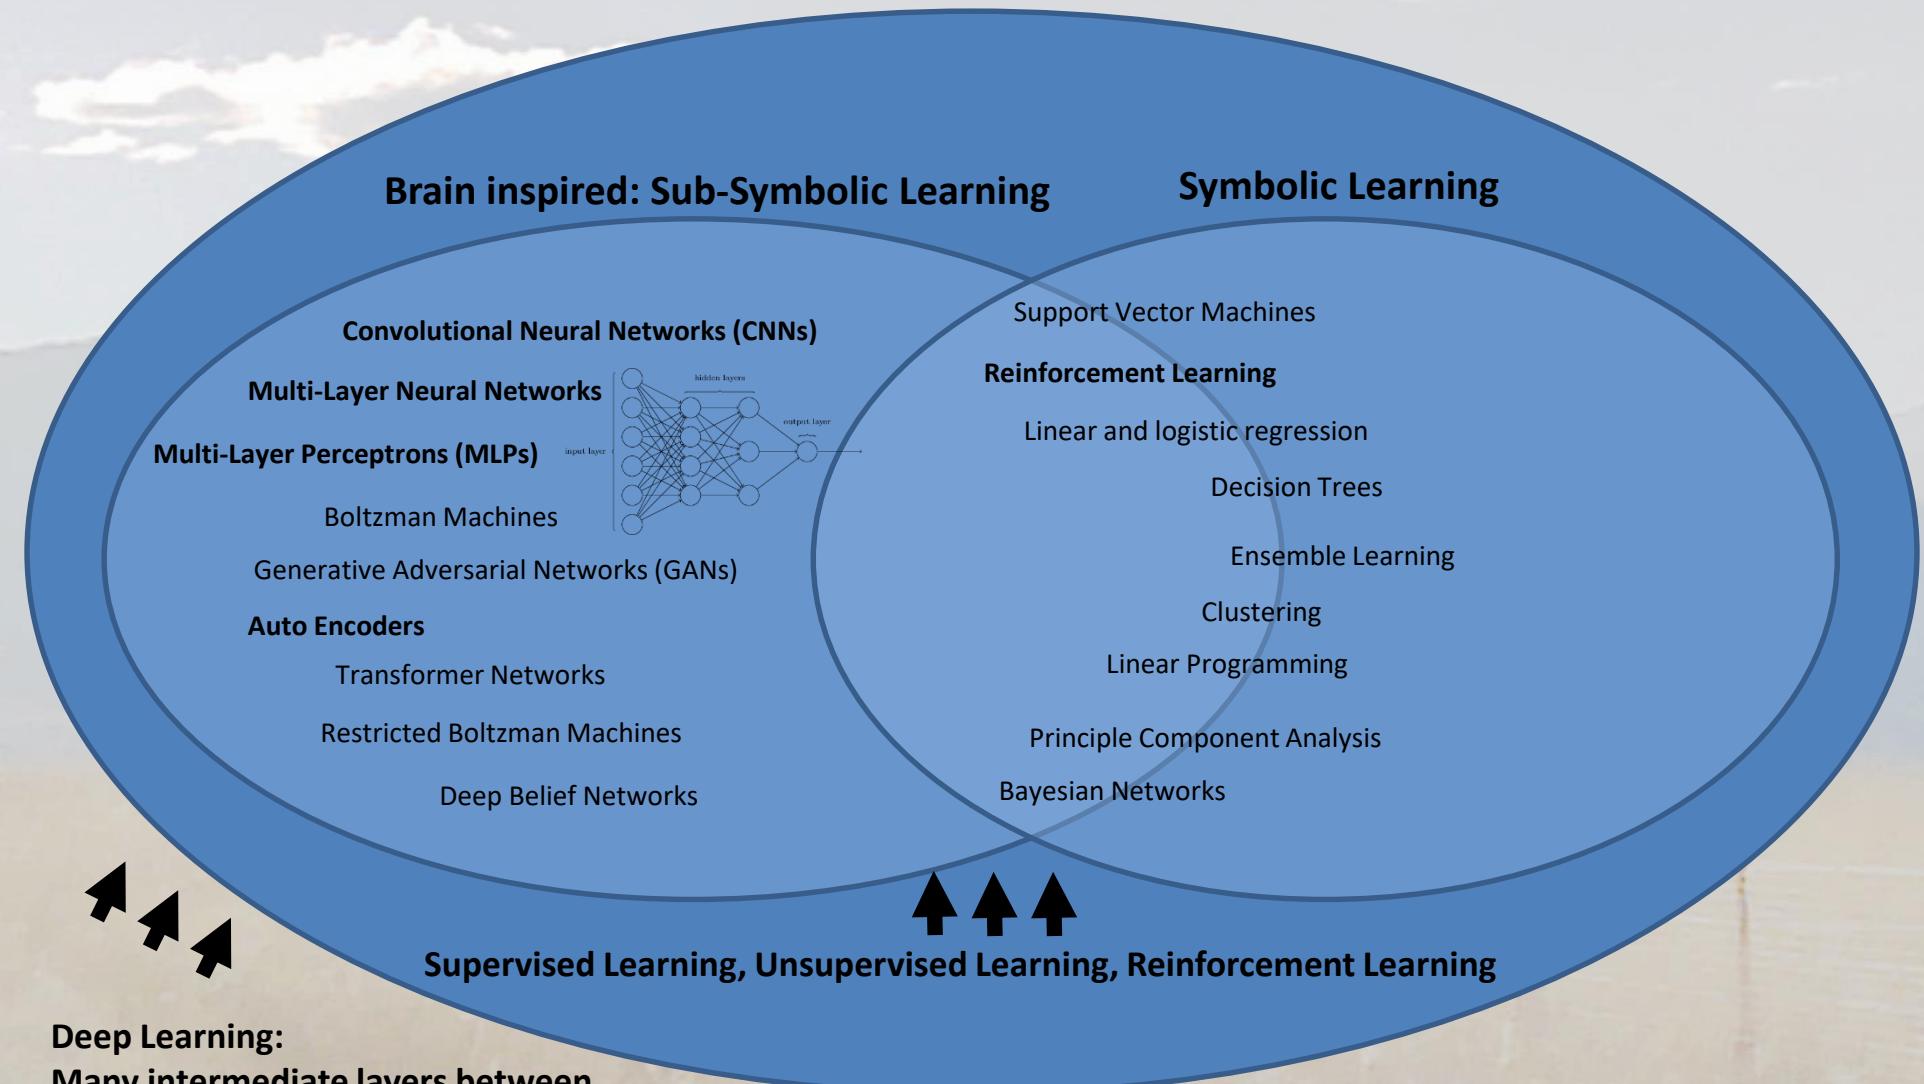
Boltzman Machines

...

- Backfed Input Cell
- Input Cell
- ▲ Noisy Input Cell
- Hidden Cell
- Probabilistic Hidden Cell
- △ Spiking Hidden Cell
- Output Cell
- Match Input Output Cell
- Recurrent Cell
- Memory Cell
- △ Different Memory Cell
- Kernel
- Convolutional or Pool



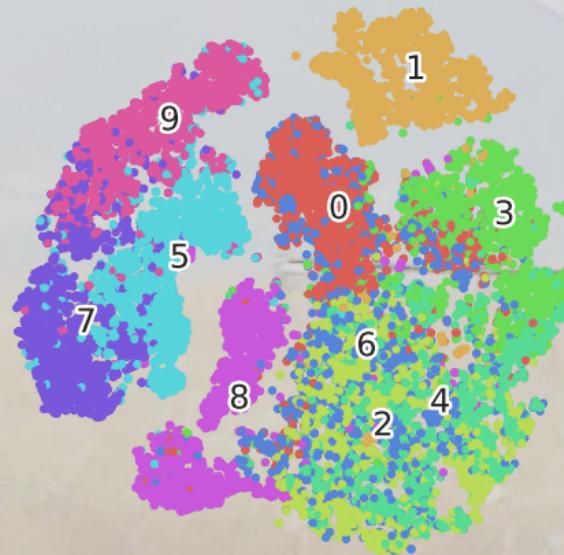
Machine Learning: we only scratched the surface



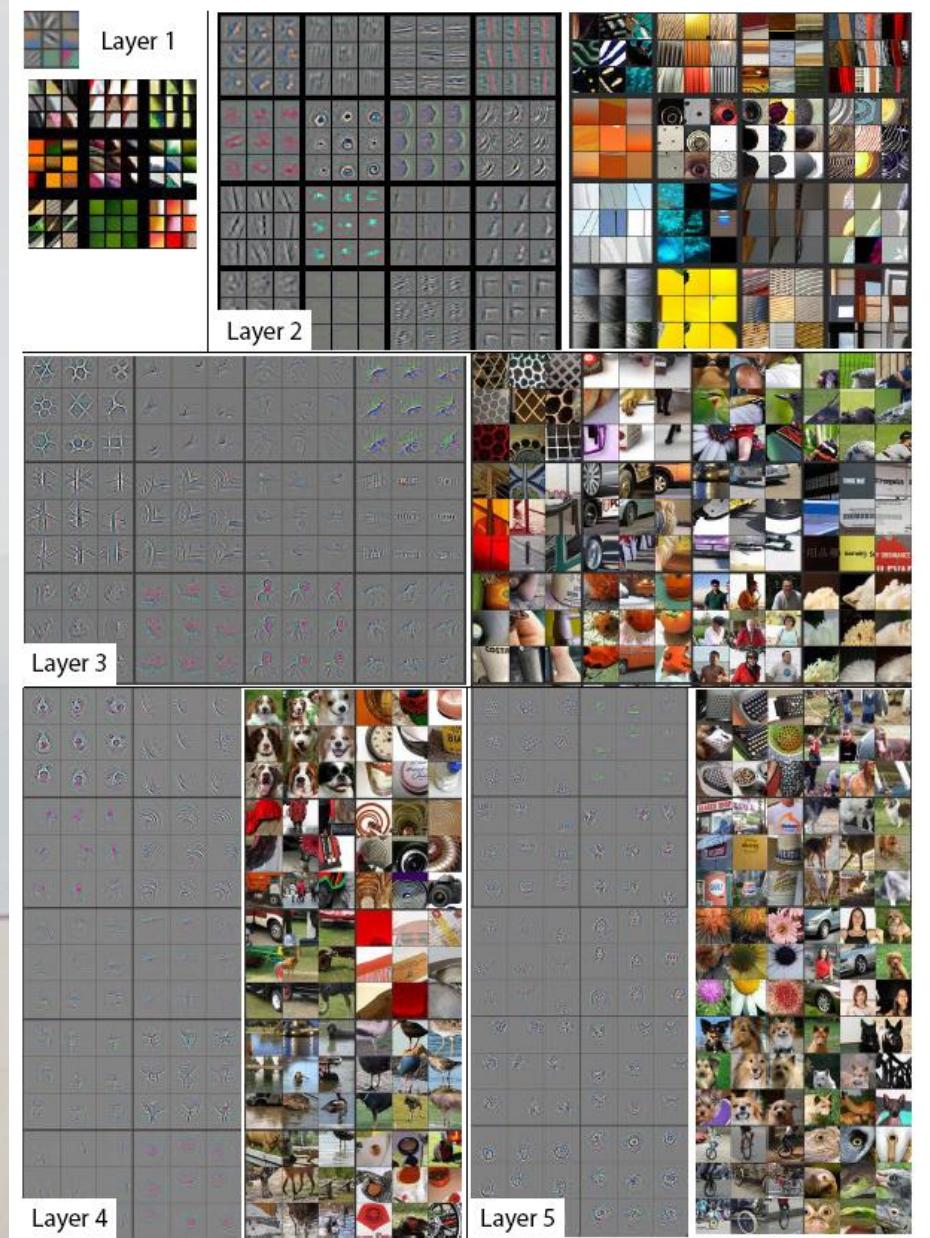
Deep Learning:
Many intermediate layers between
input and output

Limitations of neural nets

- Black box, unexplainable
- Critique: It's just associations and curve fitting
- No concept of causality



Visualization with T-SNE



And thousands of applications ...

- Advertising: Consumer preferences, movie recommendations, etc ...
- NLP Speech recognition, translation, document generation, chatbots (Alexa, Siri, ...)
- Autonomous driving
- Just type AI in your favorite search box! AI powered this, AI powered that! ☺ ☹ ☺
- Beware the hype! Not even close to passing Turing Test (imho), forget about AGI!
- <https://talktotransformer.com/> Amazing, but ...

Hype: Caveat Emptor!



Hardware and programming frameworks

Hardware

- CPU
- GPU
- USB TPU dongle
- Google Colabs, Azure, AWS, ... free time in some cases
- Edge AI

GAP8 enables AI at the very Edge

Markt&Technik Elektronik

A highly integrated MCU combining a 8 core parallel compute cluster and a single core

01 High compute at ultra-low-power
~20x better power efficiency than the state-of-the-art on content understanding applications at < 100mW

02 Agile
Ultra-fast dynamic power state transitions Wake up in 0.5ms 1uA standby current

03 Flexible
Fully programmable Wide range of accelerated algorithms

Company Proprietary

GAP8
The IoT Application Processor

Gap 8 SoC



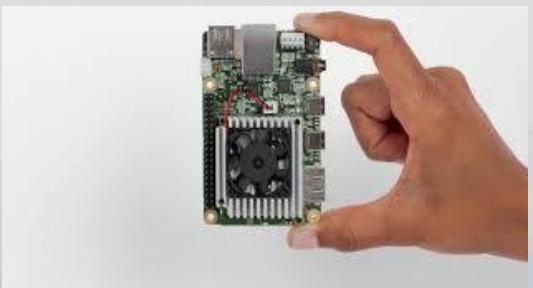
Intel Movidius



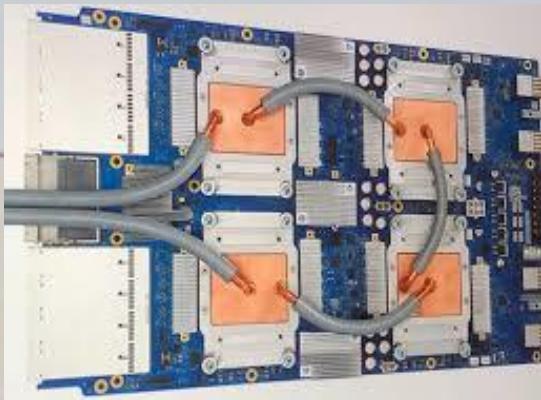
USB Google Coral



Nvidia Jetson Nano

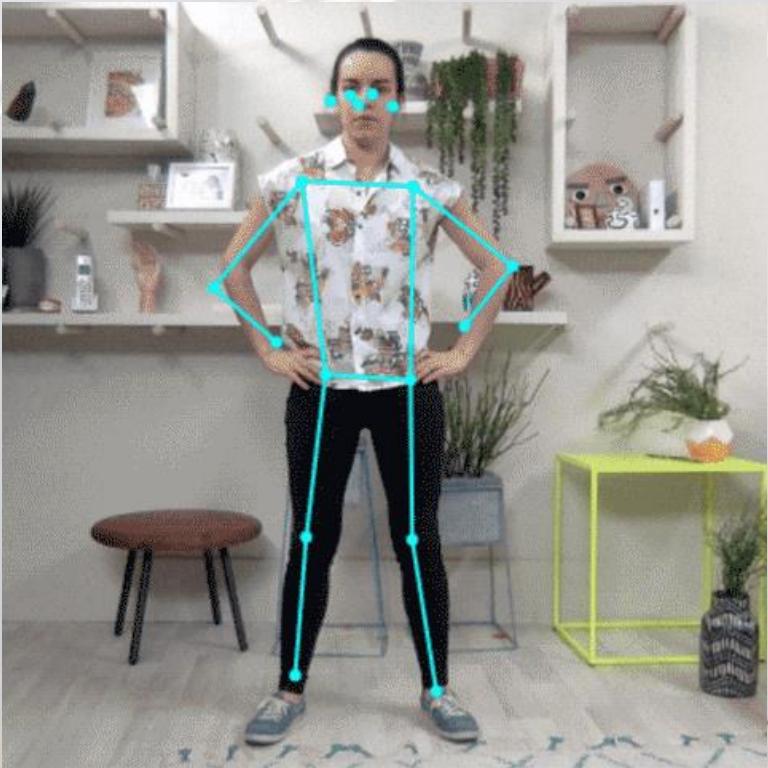


Google Coral Board

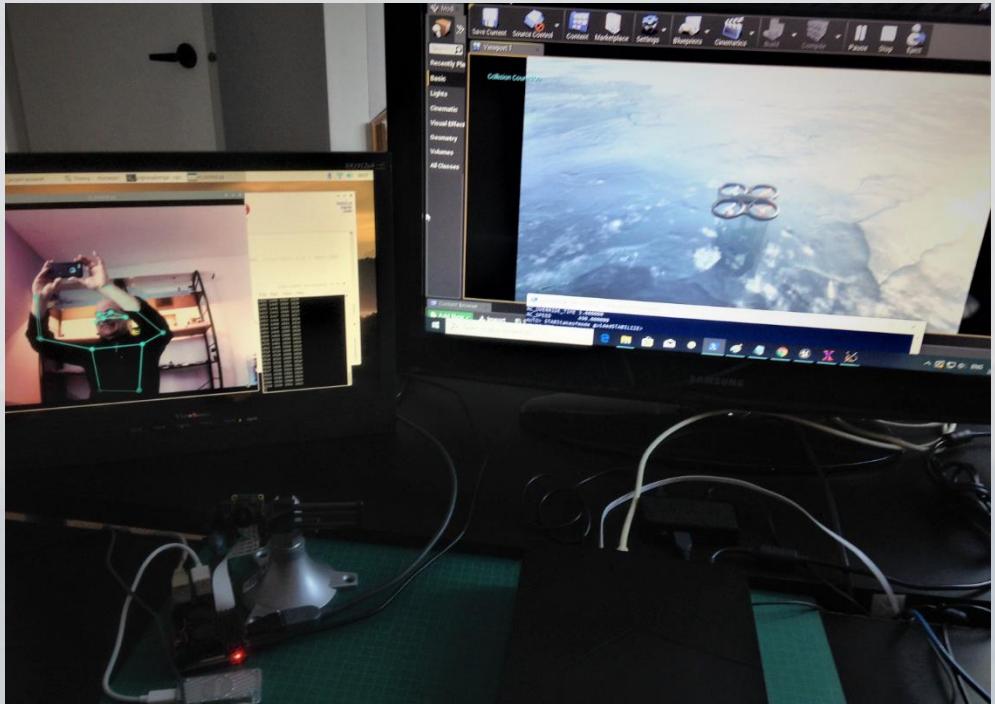


Google TPU on cloud

Hardware



PoseNet



Flying with Gestures Patrick Poirier
Pi4 + Google Coral 20FPS with posenet
<https://gitter.im/ArduPilot/VisionProjects>

Programming options

- Python, Raw C/C++/CUDA, ☹
- Keras: Tensorflow API, multibackend (tensorflow, CNTK, Theano) tf.keras:
Tensorflow only, integrated in TensorFlow
Very powerful despite simplicity
- TensorFlow (Google), PyTorch (Facebook)
both Open Source)
- HAL (GPU, CPU, TPU, X on cloud) No worries about underlying hardware.
- Rich ecosystems
- Also Tensorflow Lite, TinyML (Arduino Nano and STM32!)
- ONNX (Open Neural Network Exchange) for sharing and framework interoperability, import/export to/from cloud



Programming options, advanced custom networks:

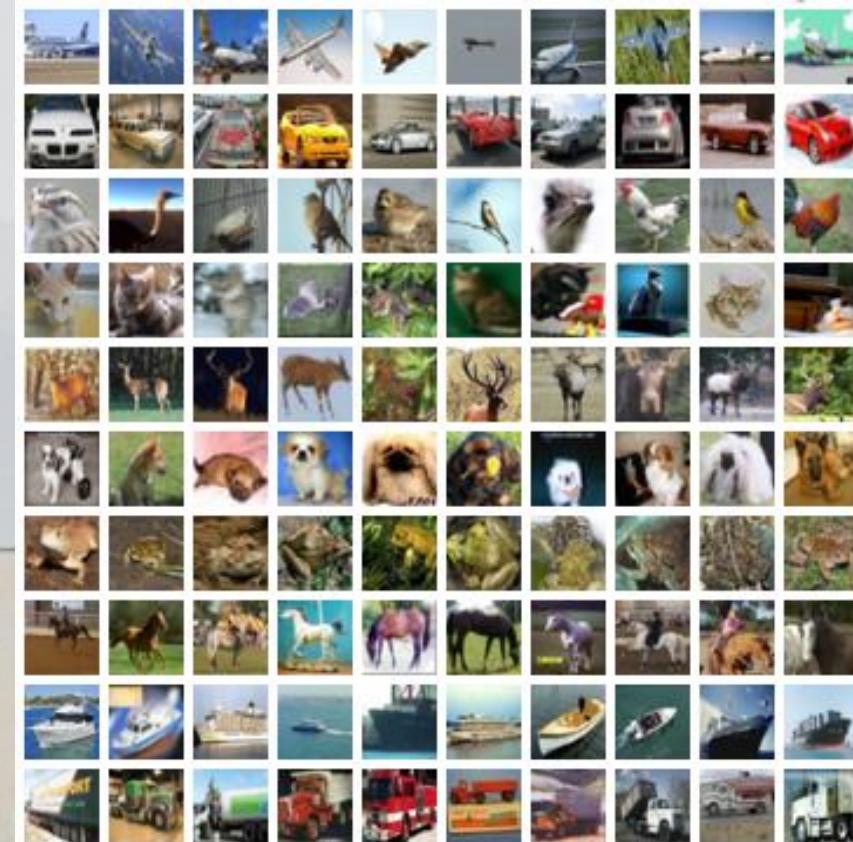
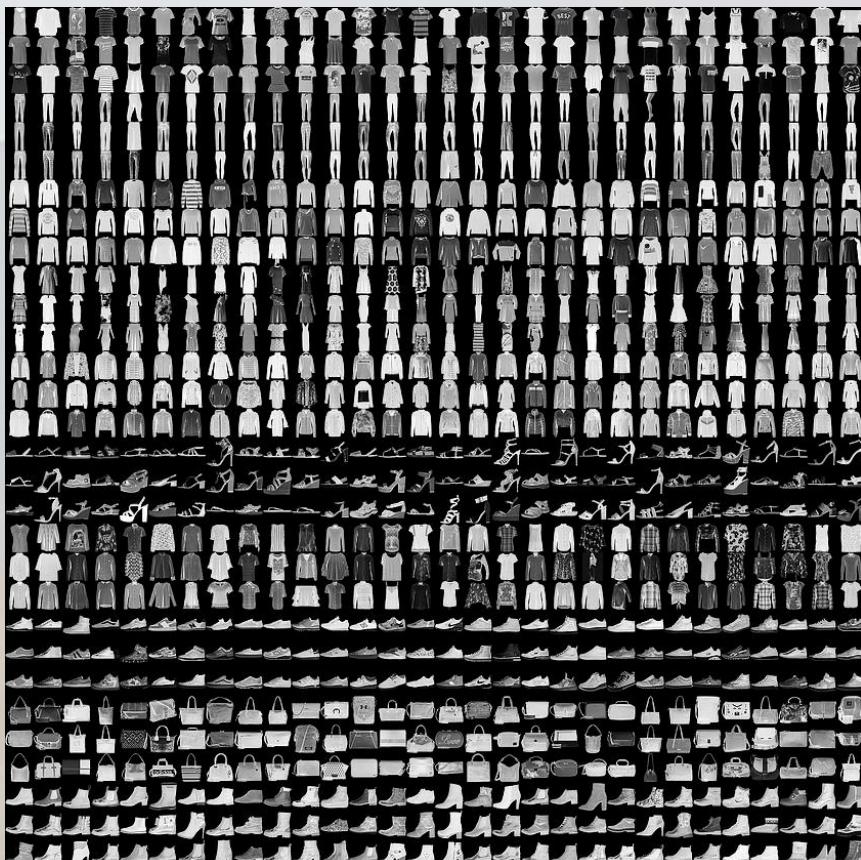
- Tensorflow Or PyTorch: <- 75% research papers (? ...)
- PyTorch and Tensorflow revolve around tensors, generalization of matrices (n,m) dimensions to (n,m,o, ...)
- Example: 1980x1080x3 channel RGB image: 3d tensor. Video: 1980x1080x(3 channel RGB) x frame#: 4D tensor
- Autodiff (TensorFlow), autograd (PyTorch): Dynamic computation graph keeping track of operations for most efficient gradient computations (Jacobians ...)



Keras: Couple simple examples to show how easy it can be to get started

FashionMnist with MLP

Cifar10 with Resnet Conv. net



Demo



AI for UAVs and UGVs

AI for flying and ground motion

- Supervised Learning or Reinforcement Learning
- Simulation or simulation+real world

Autonomous Driving, rovers

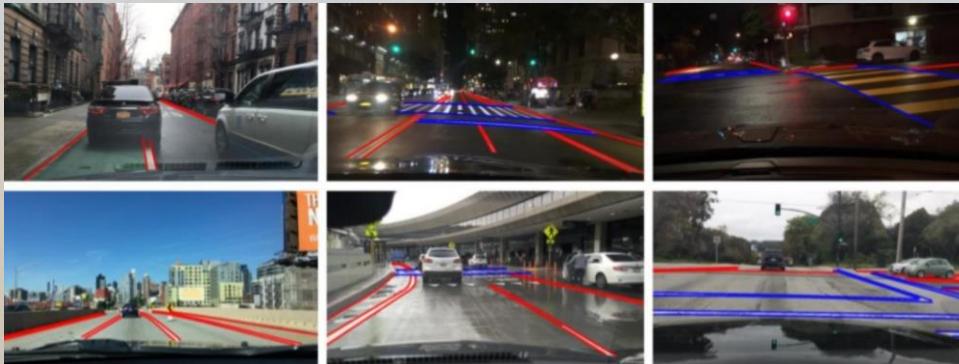
- NVIDIA: LaneNet, WaitNet, SignNet
- LightNet, ClearSightNet
- Ride in NVIDIA Self-driving Car:
- <https://www.youtube.com/watch?v=1W9q5SjaJTc>



Driving Datasets

- Berkeley Deep Drive BDD 100k
- Baidu ApolloScope Dataset
- Comma.AI Dataset
- Oxford's Robotic Car Dataset
- Cityscapes Dataset
- Kitti Dataset
- Ford Campus Vision And Lidar Dataset
- Motion-based Segmentation And Recognition Dataset
- TuSimple Dataset
- CMU Visual Localization Dataset
- CCSAD Dataset
- Kul Belgium Traffic Sign Dataset
- MIT Age Lab Dataset
- Lisa: Intelligent & Safe Automobiles, UCSD Datasets
- Udacity Challenge Datasets
- NCLT Datasets
- DIPLECS Autonomous Driving Datasets
- Velodyne SLAM Dataset
- Daimler Urban Segmentation Dataset
- The Uah-driveset
- DAVIS Driving Dataset 2017 (DDD17)
- Berkeley Deepdrive (BDD) Driving Model
- MIT-AVT: Autonomous Vehicle Technology

<https://sites.google.com/site/yorkyuhuang/>
Kang et al. 2019



Berkeley dataset, annotated



Udacity dataset

End to end supervised learning

- ALVINN, CMU, 1989

What's Hidden in the Hidden Layers?

*The contents can be easy to find with a geometrical problem,
but the hidden layers have yet to give up all their secrets*

David S. Touretzky and Dean A. Pomerleau

AUGUST 1989 • B Y T E 231

tions, we fed the network road images taken under a wide variety of viewing angles and lighting conditions. It would be impractical to try to collect thousands of real road images for such a data set. Instead, we developed a synthetic road-image generator that can create as many training examples as we need.

To train the network, 1200 simulated road images are presented 40 times each, while the weights are adjusted using the back-propagation learning algorithm. This takes about 30 minutes on Carnegie Mellon's Warp systolic-array supercomputer. (This machine was designed at Carnegie Mellon and is built by General Electric. It has a peak rate of 100 million floating-point operations per second and can compute weight adjustments for back-propagation networks at a rate of 20 million connections per second.)

Once it is trained, ALVINN can accurately drive the NAVLAB vehicle at about 3½ miles per hour along a path through a wooded area adjoining the Carnegie Mellon campus, under a variety of weather and lighting conditions. This speed is nearly twice as fast as that achieved by non-neural-network algorithms running on the same vehicle. Part of the reason for this is that the forward pass of a back-propagation network can be computed quickly. It takes about 200

milliseconds on the Sun-3/160 workstation installed on the NAVLAB.

The hidden-layer representations ALVINN develops are interesting. When trained on roads of a fixed width, the net-

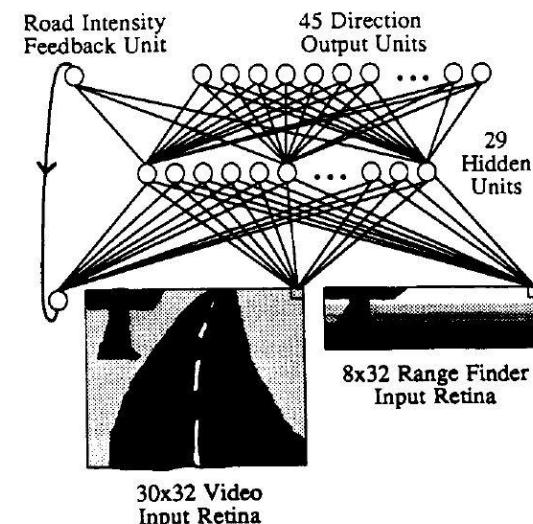
work chooses a representation in which hidden units act as detectors for complete roads at various positions and orientations. When trained on roads of variable

continued



Photo 1: The NAVLAB autonomous navigation test-bed vehicle and the road used for trial runs.

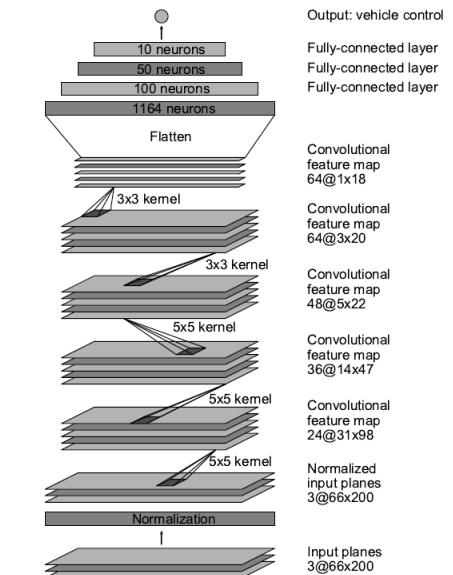
ALVINN Architecture



Pomerleau, 1989

End to end learning

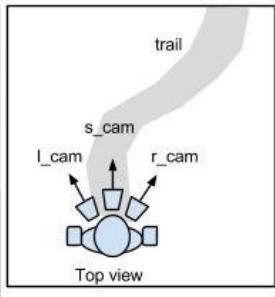
- Nvidia, 2016. Input: road images, output: steering



Nvidia, 2016

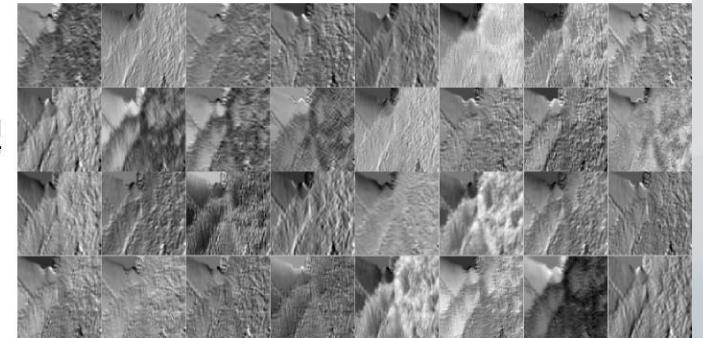
Drone Forest Trail Autonomous Navigation, 2016

<http://people.idsia.ch/~giusti/forest/web/>

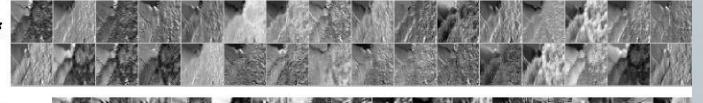


- Training examples acquired with 3 Gopros, 8 hours of video, 7km of trails, 20k+ images
- Convnet: 10 layers, 150k Weights, 500k Neurons, 57 Million Connections
- Parrot drone with net running on laptop via wifi
- Experiments also made with onboard Odroid-U3, 15fps achieved

L0 - Input layer: 3 maps of 101x101



L1 - Convolutional Layer: 32 maps of 98x98 neurons. Filter: 4x4



L2 - MaxPooling Layer: 32 maps of 49x49 neurons. Kernel 2x2



L3 - Convolutional Layer: 32 maps of 46x46. Filter 4x4



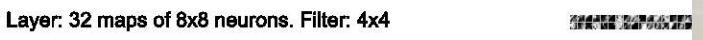
L4 - MaxPooling Layer: 32 maps of 23x23. Kernel: 2x2



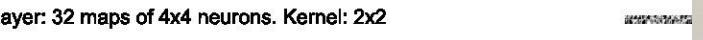
L5 - Convolutional Layer: 32 maps of 20x20. Filter: 4x4



L6 - MaxPooling Layer: 32 maps of 10x10 neurons. Kernel: 2x2



L7 - Convolutional Layer: 32 maps of 8x8 neurons. Filter: 4x4



L8 - MaxPooling Layer: 32 maps of 4x4 neurons. Kernel: 2x2



L9 - Fully Connected Layer: 200 neurons



L10 - Output Layer: 3 neurons



Success (top) and failures (bottom)



Olivier J. Brousse 3/26/2020

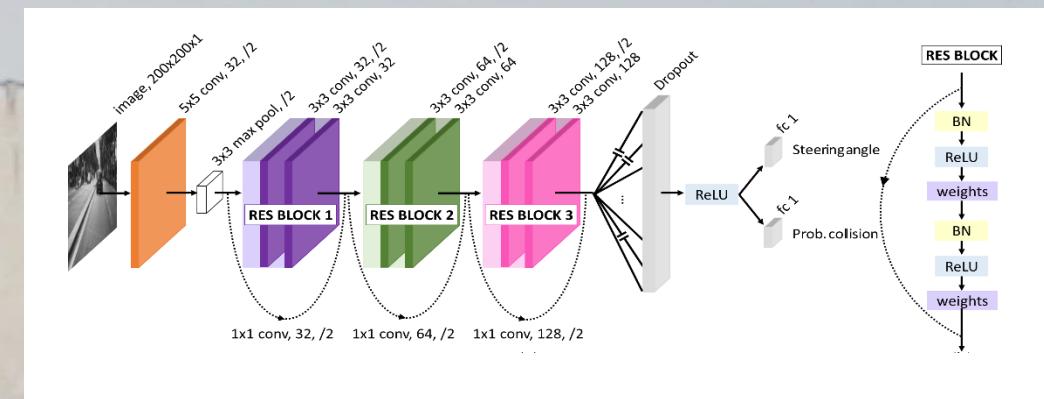
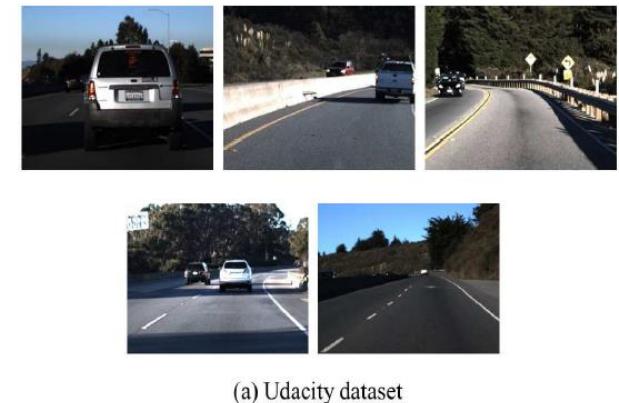
ARDUPILOT

3DVistas
Drones for Research and Industry

DroNet, Learning to fly by driving, 2018

http://rpg.ifi.uzh.ch/docs/RAL18_Loquercio.pdf

- ResNet architecture, trained with Keras/Tensorflow
- Split output: Steering Angle, Collision probability
- Two training datasets: Udacity car driving, 70k images for steering angle prediction; Gopro on bicycle handlebar, 32k images, manually annotated, for collision probability.
- Parrot Bebop; Velocity controlled from Core I7 laptop via wifi



DroNet, Learning to fly by driving, 2018

<https://www.youtube.com/watch?v=ow7aw9H4BcA>

Code, datasets open sourced.

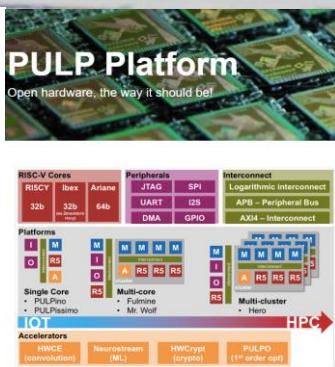
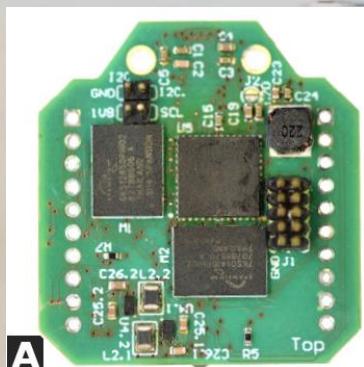
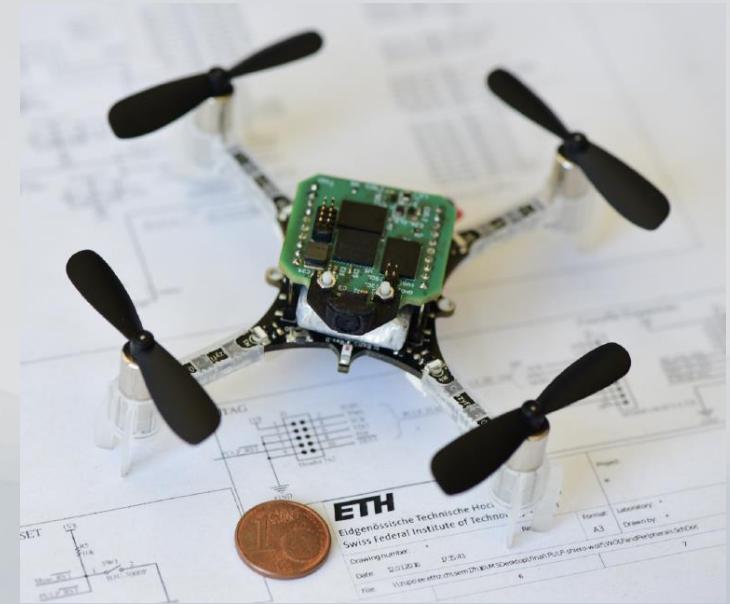


Pulp-DroNet, 2019

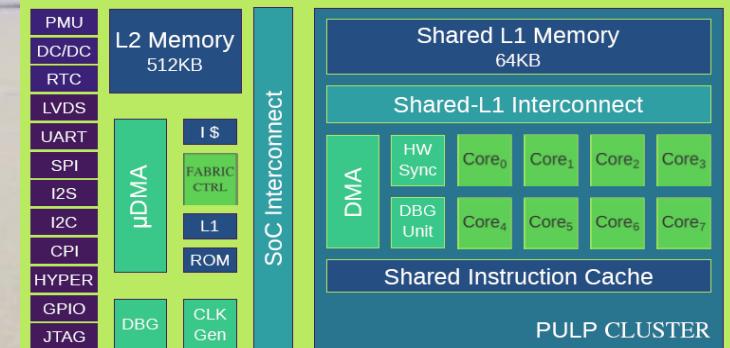
<https://github.com/pulp-platform/pulp-dronet/>

<https://www.youtube.com/watch?v=57Vv5cSvnaA>

- Crazyflie 2.0 with Pulp-Shield open hardware
- 27g, power draw < 300mw, 6 Fps visual processing
- Gap 8 SoC, 8 core, Greenwave Technologies
- Same network architecture as before, tweaked:
Quantization (float32 to Fixed16), receptive field
from max pooling layers size from 3x3 to 2x2



Gap 8 SoC



Pulp-DroNet, 2019

<https://github.com/pulp-platform/pulp-dronet/>

<https://www.youtube.com/watch?v=57Vy5cSvnaA>



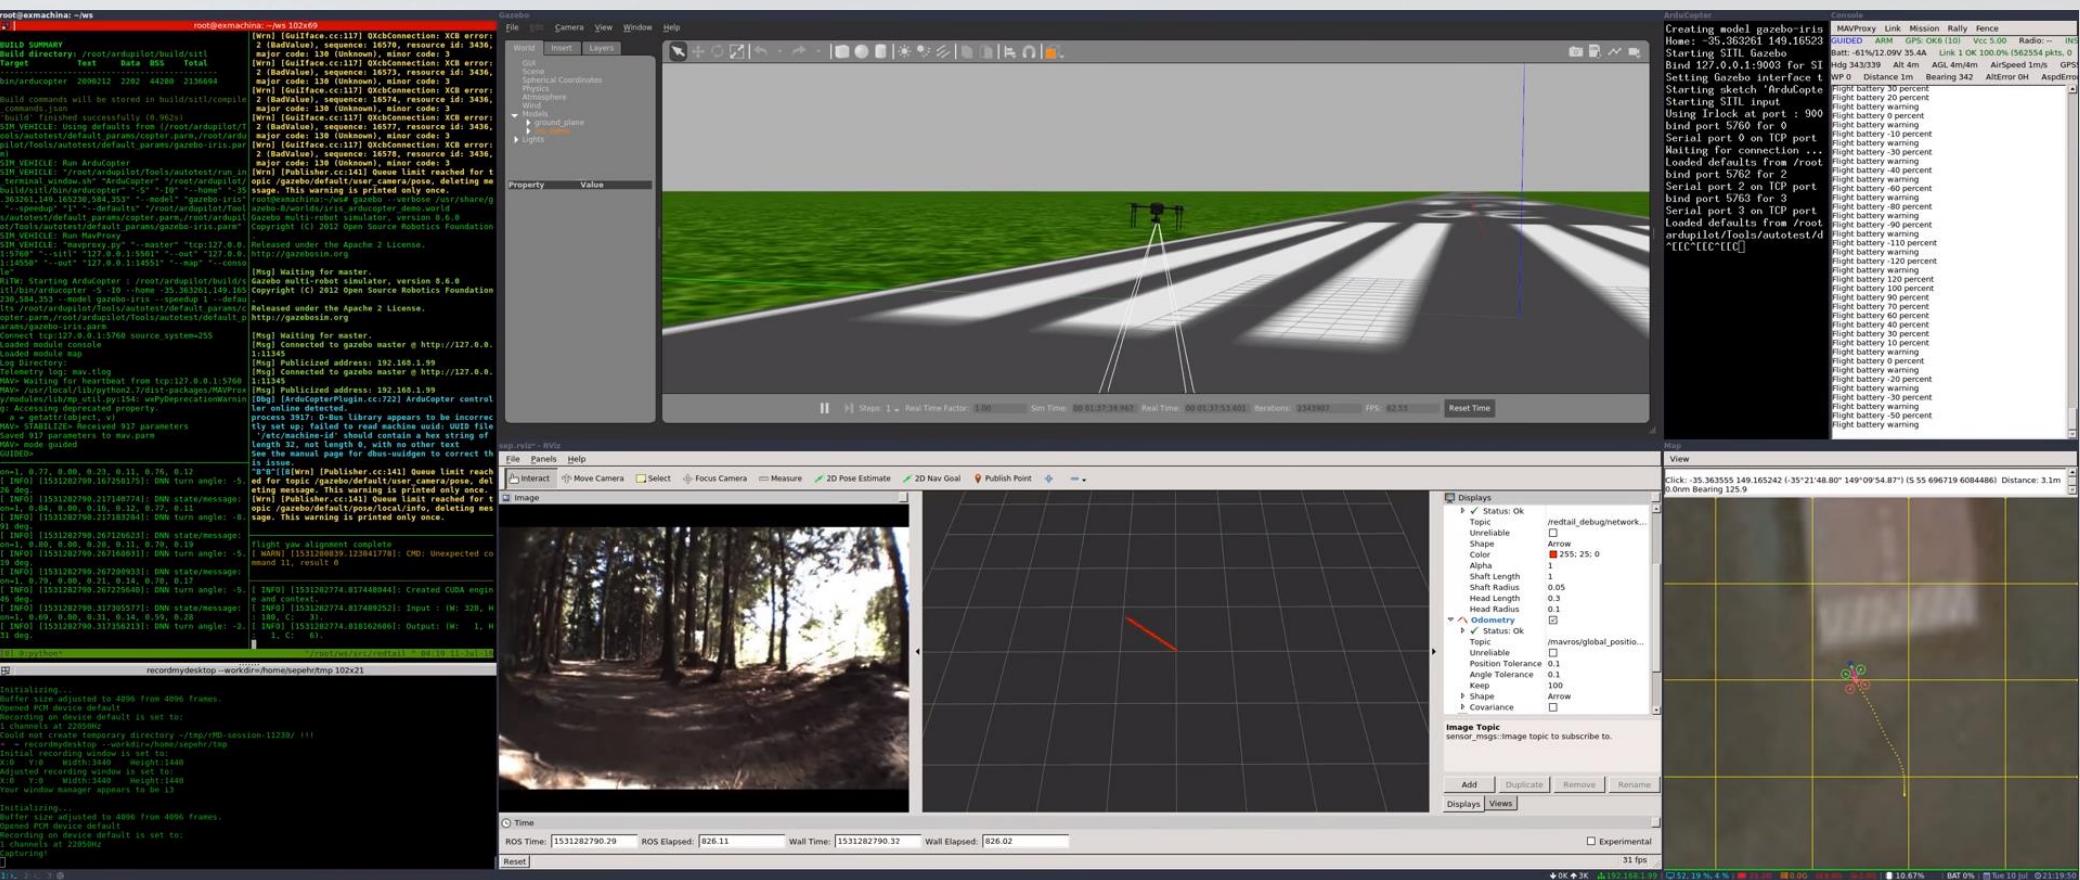
Copyright 2019 © ETH zürich

Olivier J. Brousse 3/26/2020

ARDUPILOT

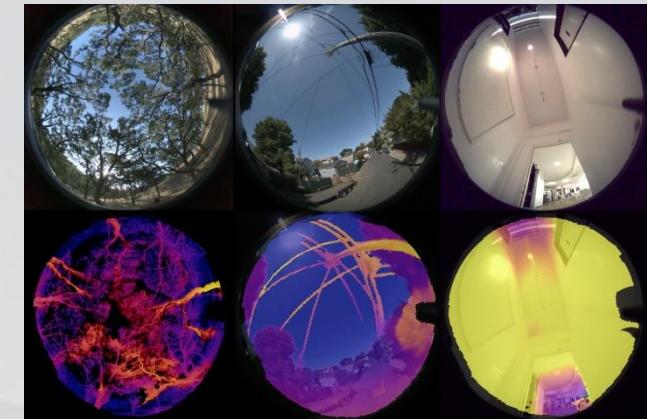
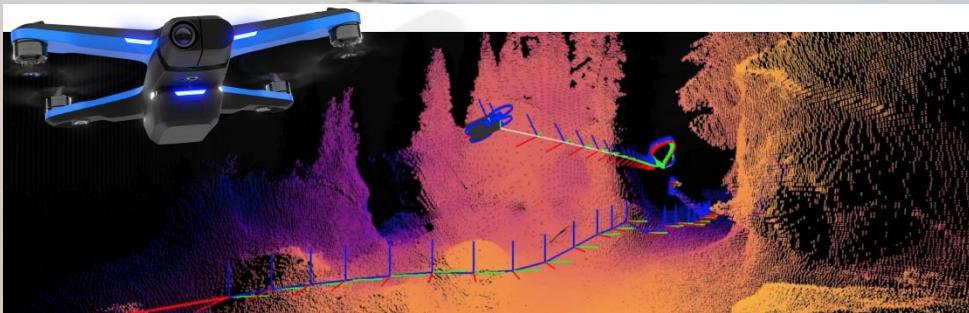
3DVistas
Drones for Research and Industry

Nvidia Red Tail, GSOC 2018 Ardupilot port



See also: SKYDIO

- Posts on Medium:
 - “Inside the mind of the Skydio 2”
 - ” Deep Neural Pilot on Skydio 2”
- CNN for 3D Map from redundant stereo pairs?
- CNN for Object tracking?
- --- ?

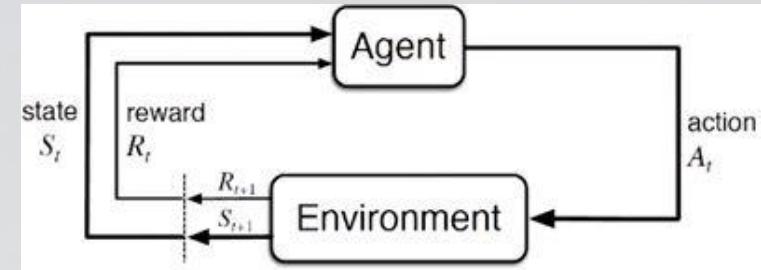


The background of the slide features a wide-angle photograph of a natural landscape. In the foreground, there is a field with some sparse vegetation and a few vertical poles, possibly remnants of a fence or marker posts. Beyond the field, a range of mountains is visible, their peaks obscured by a hazy, light-colored atmosphere. The sky above is a clear, pale blue with scattered, wispy white clouds.

(Deep) Reinforcement Learning

Reinforcement Learning

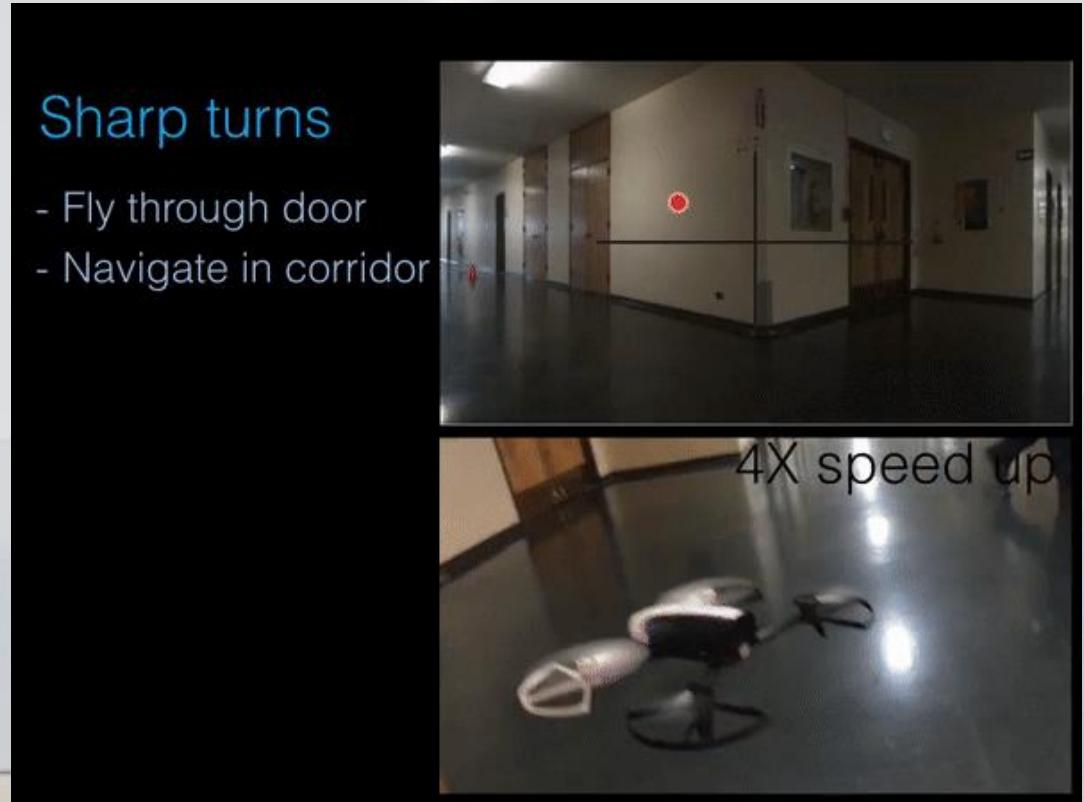
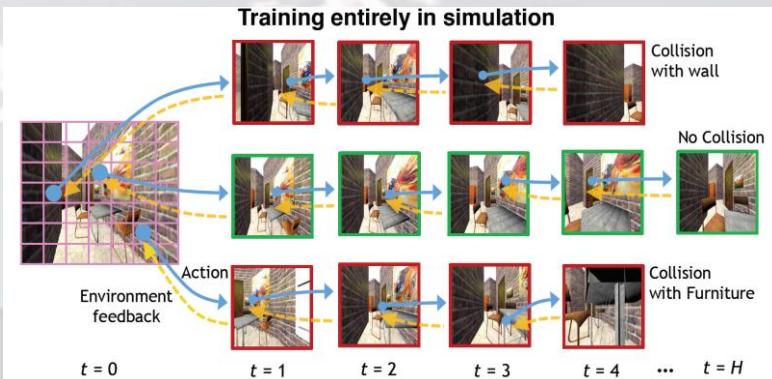
- Explore space: Fly randomly
- Crash, Hit/Avoid obstacle: Penalty/Reward
- Rinse and repeat, 1000's of times and learn optimal policy



- AirLearning
- Built on top of AirSim and OpenGym
- <https://github.com/harvard-edge/airlearning>



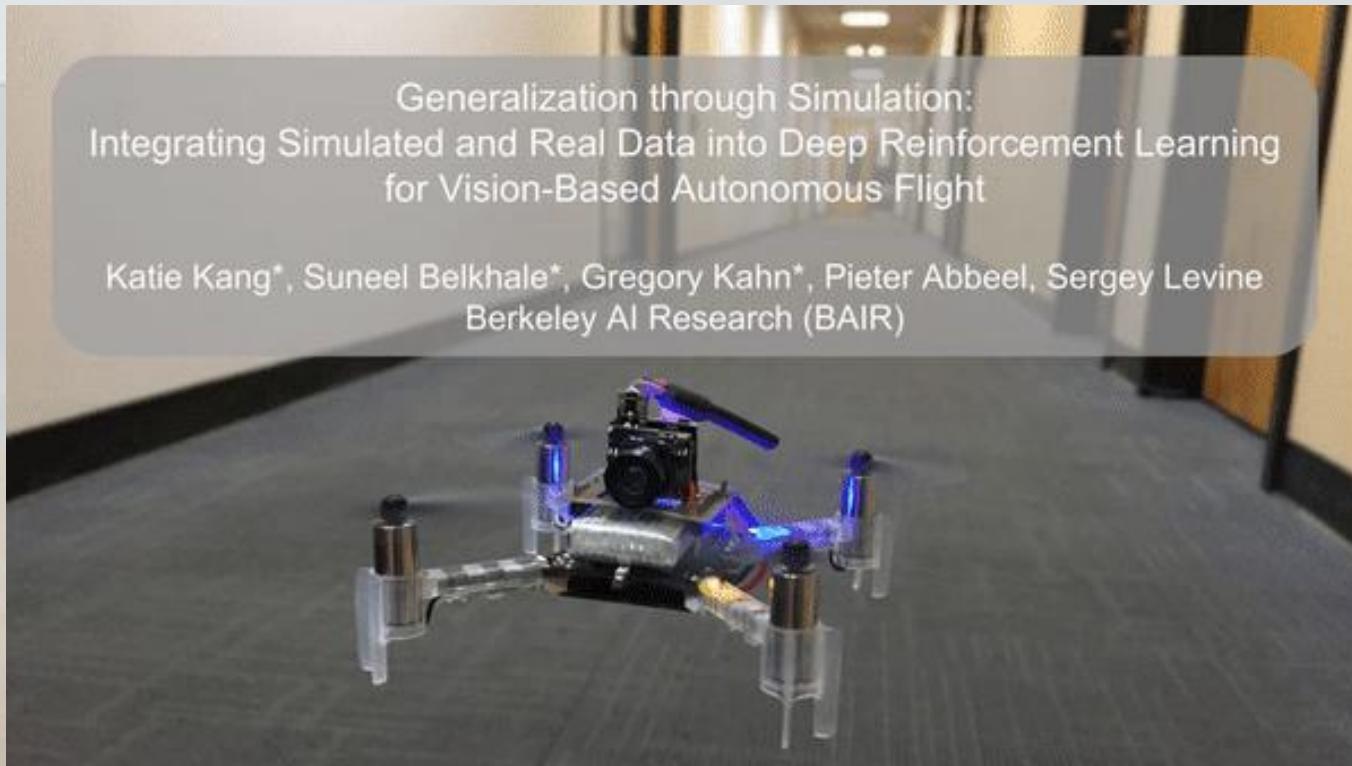
CAD²RL: “Real Single-Image Flight Without a Single Real Image”, 2017



- Deep RL (VGG16)
Training with Blender
generated environment images
- VGG-16 Net. Randomize examples
with different textures, lighting, objects,
etc ... to be able to generalize to
real world

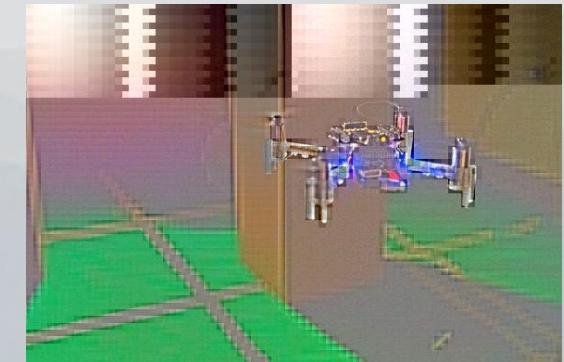
DeepRL (Cont) 2019

- <https://github.com/gkahn13/GTS>
- Generalization through Simulation: Integrating Simulated and Real Data into Deep Reinforcement Learning for Vision-Based Autonomous Flight, 2019
- BitCraze Crazyflie



DeepRL (cont). See also:

- “Autonomous Navigation via Deep Reinforcement Learning for Resource Constraint Edge Nodes Using Transfer Learning”, Oct. 2019 (DJI Tello)
https://github.com/aqeelanwar/DRLwithTL_real
- Learning to Seek: Autonomous Source Seeking with Deep Reinforcement Learning Onboard a NanoDrone Microcontroller” Using Transfer Learning”, Sep 2019 (CrazyFlye) ”
<https://github.com/harvard-edge/source-seeking>
- “Learning to fly by crashing”, 2017,
<https://youtu.be/u151hJaGKUo>





Autonomous Drone Racing

Olivier J. Brousse 3/26/2020

ARDUPILOT

3DVistas
Drones for Research and Industry

Autonomous Drone Racing

- Significant interest by the research community given input simplicity yet problem complexity
- Difficult for SLAM and visual odometry
- Challenges:
 - Dynamic environment
 - Drift with visual odometry
 - Speed
 - Limited onboard compute resources

Drone Racing: Competitions

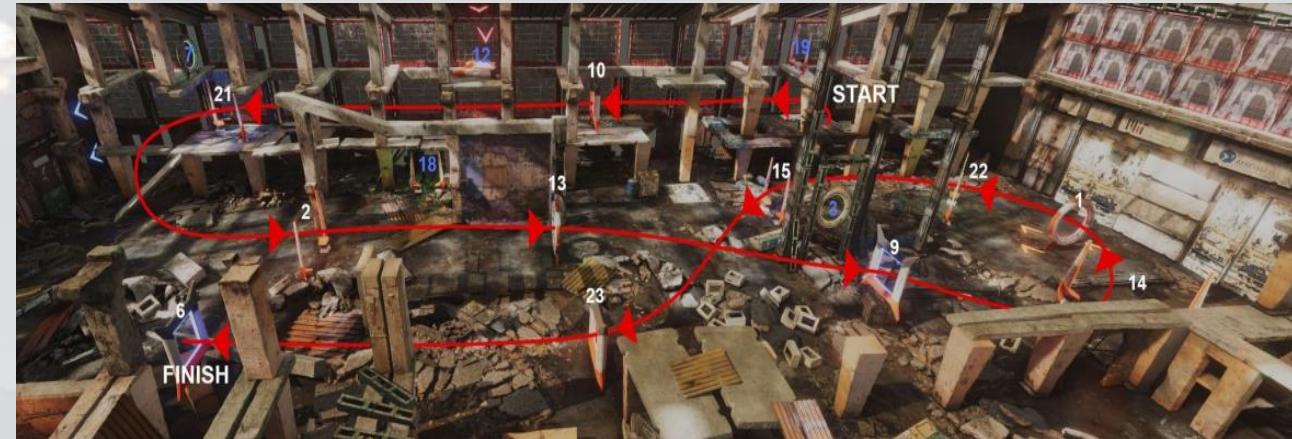
AlphaPilot – Lockheed Martin AI Drone Racing Innovation Challenge

AlphaPilot is the first large-scale open innovation challenge of its kind focused on advancing artificial intelligence (AI) and autonomy.

Data Science Drones Technology

Stage:
2020 Season Coming Soon!

Prize:
\$2,250,000



FlightGoggles <https://github.com/mit-fast/FlightGoggles>

- AlphaPilot. Pre-qualification with code sent to Lockheed Martin running on FlightGoggles simulator
- IROS ADR (Autonomous Drone Race) International Conference on Intelligent Robots and Systems



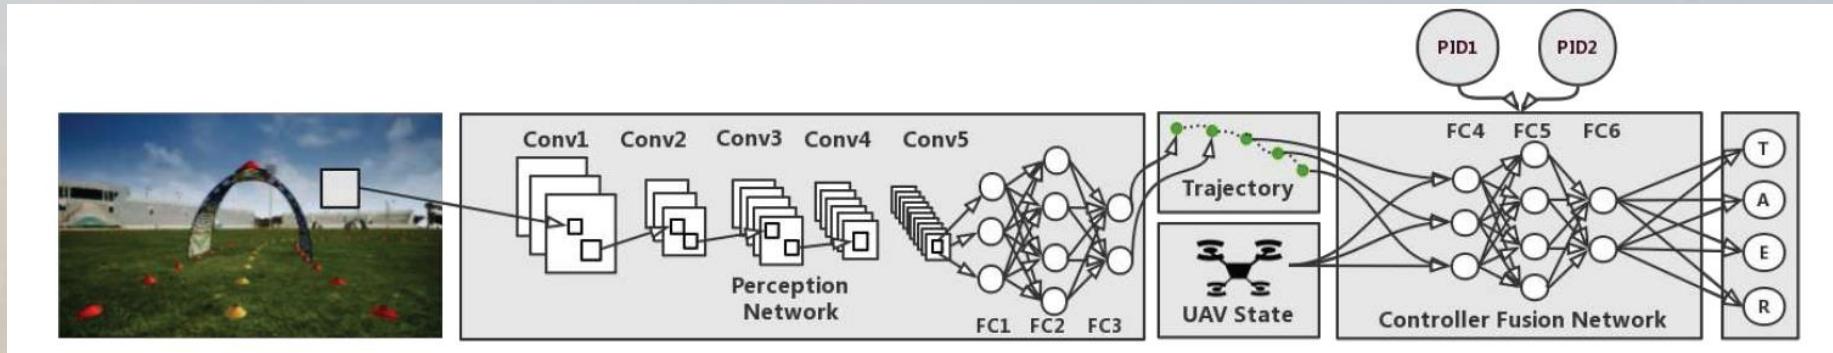
Drone Competitions

- Competition at NeuRIPS 2019: Game of Drones
- Based on AirSim
- <https://github.com/microsoft/AirSim-NeurIPS2019-Drone-Racing>
- <https://microsoft.github.io/AirSim-NeurIPS2019-Drone-Racing/>
- (Presentation, overview of participants approach, participants reports)



Teaching UAVs to race (KAUST)

- “Teaching UAVs to Race: End-to-End Regression of Agile Controls in Simulation, 2018”
- “Learning a Controller Fusion Network by Online Trajectory Filtering for Vision-based UAV Racing”, 2019
<https://www.youtube.com/watch?v=hGKIE5X9Z5U>,
<https://www.youtube.com/watch?v=9cbjOmKwbUY>
- End to end via imitation learning
- Training and running on photo-realistic simulation, SIM4cv on top of UnReal, <https://sim4cv.org/>

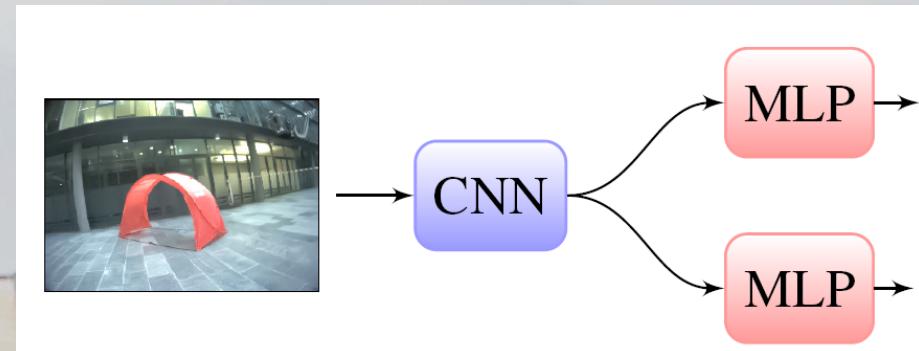


Teaching UAVs to race (KAUST)



ETH Zurich http://rpg.ifi.uzh.ch/research_drone_racing.html

- “Beauty and the Beast: Optimal Methods Meet Learning for Drone”, 3/2019
- “Deep Drone Racing: from Simulation to Reality with Domain Randomization”, 11/2019
- Predicted waypoints in local body frame with CNN then fed to planner and tracker



Training samples:
Collected via measurement
with manual flying

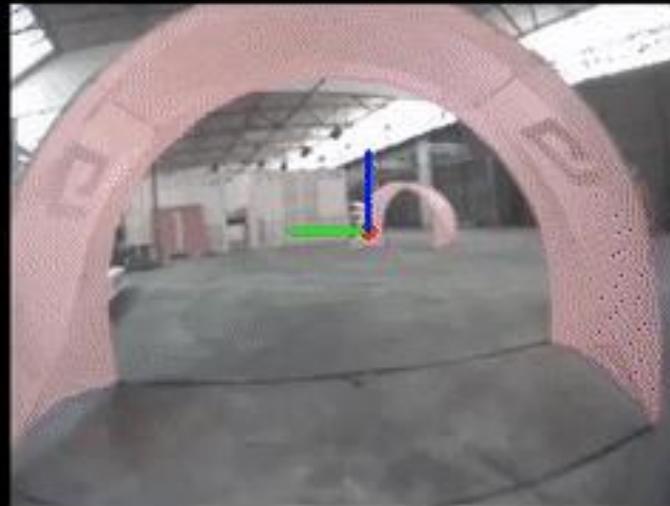
Output:
Mean and Variance of
distribution describing next
gate pose estimate
distribution

ETH Zurich http://rpg.ifi.uzh.ch/research_drone_racing.html

3.5 m/s

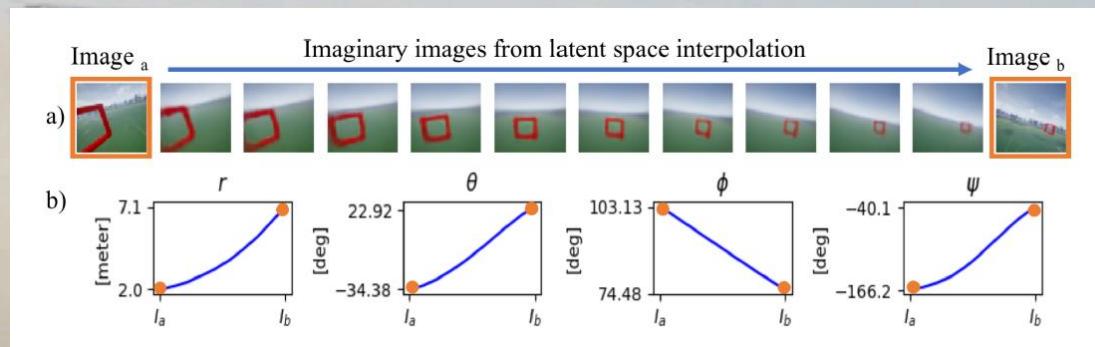


3rd person view



1st person view
(predicted gate pose overlaid)

- “Learning Controls Using Cross-Modal Representations: Bridging Simulation and Reality for Drone Racing”, 9/2019
- Neural net used: variational autoencoder. Training of VA forces hidden representation to be as close as normally distributed as possible. => can reconstruct from arbitrary samples([Variational Autoencoders Tutorial](#)).
- Training Input/output: Cross modal: RGB image (Airsim), (body frame spherical coordinates and yaw), Dronet architecture, Resnet used.
- Training sample generation: Use planner, one gate ahead horizon, 14k data points



See also: “Autonomous drone race: A computationally efficient vision-based navigation and control strategy”, Nov 2018

<https://arxiv.org/pdf/1809.05958.pdf>

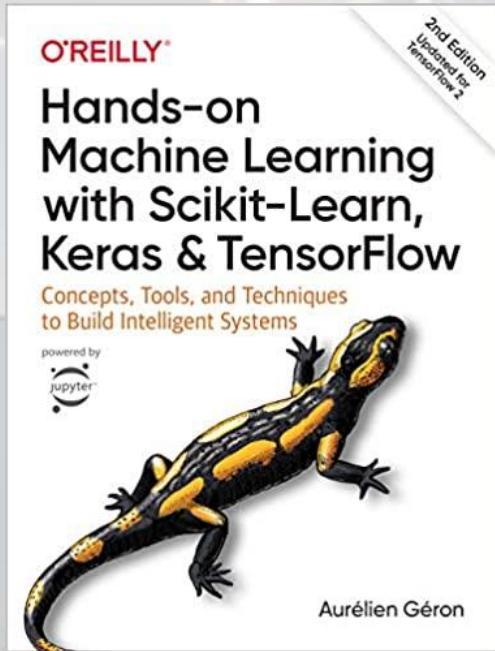
- TU-Delft, Delft University of Technology, The Netherlands
- Parrot Bebop with Paparazzi autopilot (Cortex A9, “dual core only”)
- Light-weight, **not CNN**, gate detection vision algorithm
- Beat ETH and won Alphapilot 2019 (\$1 million)
- So there! ☺

Links, Survey Papers

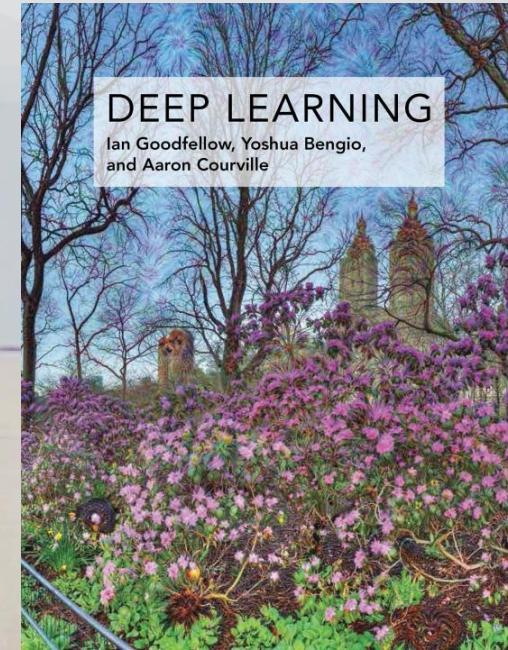
- “A survey of deep learning techniques for autonomous driving”, Oct. 2019, <https://arxiv.org/pdf/1910.07738.pdf>
- “A Review on IoT Deep Learning UAV Systems for Autonomous Obstacle Detection and Collision Avoidance”, Sep 2019, <https://www.mdpi.com/2072-4292/11/18/2144>
- “A Review of Deep Learning Methods and Applications for Unmanned Aerial Vehicles”, 2017, <https://www.hindawi.com/journals/js/2017/3296874/>
- Also of interest: Papers with code, <https://paperswithcode.com/sota>, if you want to check out the latest and greatest.

Recommended books

- Applied:



Theoretical deep dive,
bible status:



The background of the slide is a photograph of a rural landscape. In the foreground, there is a field with a wire fence made of wooden posts and horizontal wires. Beyond the fence, there are several layers of mountains. The sky is filled with large, white, fluffy clouds.

Thank you!